

An Interactive Java Technical Review Simulator

Java AI Interview Simulator

Viva & Examination Preparation Guide

RITHIKA S (2104251040779) · **SHWETHAA JV** (2104251050318)

Supervisor: **Dr. R. Kavitha**, M.E., Ph.D., Associate Professor · HoD: Dr. S. Pavithra, M.E., Ph.D.
B.E. Computer Science and Engineering · Academic Year 2026–2027

What this guide contains

- A 2-minute project pitch and key numbers to quote
- System architecture, class diagrams, flowcharts and a sequence diagram
- Every design pattern used — what the pattern is, where it is in our code, and why
- OOP and SOLID principles mapped to real classes
- A class-by-class walkthrough of the codebase
- 60+ project-specific viva questions with model answers
- Modern Java features used in the code, explained
- 120+ core-Java questions (OOP, Collections, Strings, Exceptions, JVM, Multithreading, Java 8+, JUnit, Maven)
- A chapter-by-chapter walkthrough of the PBL report, with questions for each team member
- Known limitations and “tricky” questions — with honest, confident answers
- A checklist of errors in the report to fix before submission

Based on the source code (package `com.interviewsim`), the A3 poster and the PBL report · October 2026

Contents

1	Project at a glance & 2-minute pitch	13	Core Java: OOP
2	Technology stack & build	14	Core Java: Collections
3	System architecture	15	Core Java: Strings
4	Class diagrams	16	Core Java: Exceptions
5	Flow diagrams	17	Core Java: JVM & memory
6	Design patterns used	18	Core Java: Multithreading
7	OOP pillars in the project	19	Java 8+ features
8	SOLID principles in the project	20	Testing (JUnit 5), Maven & Git
9	Class-by-class walkthrough	21	The PBL report: walkthrough & questions
10	Key algorithms explained	22	Limitations, tricky questions & improvements
11	Project viva questions & answers	23	Last-minute revision sheet
12	Java language features used in the code	A	Appendix: report errors to fix

How to use this guide. Read chapters 1–3 until you can draw the architecture from memory. Learn the four main patterns in chapter 6 (Strategy, Factory, Decorator, Observer) well — examiners almost always ask “which design patterns did you use and why?”. Chapter 21 follows the report chapter by chapter, because examiners usually ask questions with the report open in front of them. Chapter 22 prepares you for the uncomfortable questions; being able to name your own project’s weaknesses and how you would fix them scores very well in a viva.

1. Project at a glance

23

questions in questions.json

6

core-Java topics (enum Topic)

34/34

JUnit 5 tests pass (11 classes)

33

source classes / interfaces in 6 packages

1.1 The 2-minute pitch (memorise this)

“Our project is **An Interactive Java Technical Review Simulator** — the app itself is called the Java AI Interview Simulator. It is a console-based placement-preparation tool written in core Java 17. Campus placements almost always have a core-Java technical round, but students usually practise with static notes that give no feedback. Our application registers a candidate, then runs a **5-question mock interview**. Questions come from a **JSON question bank** of 23 questions across six topics — OOP, Collections, Strings, Exceptions, JVM and Multithreading — loaded with **Gson**. The user picks a **question-setting strategy**: Random, Balanced (round-robin across topics) or Topic-focused.

Each typed answer is scored by a **pluggable evaluator**. If an AI provider is configured — Claude, Gemini, or a local Ollama model — the answer is graded by an LLM; otherwise, or if the AI fails mid-session, a **FallbackEvaluator** automatically switches to an **offline keyword-based evaluator**. So the app always works without internet.

At the end it prints a **placement-readiness report**: overall score, topic-wise breakdown with bar charts, a verdict — Ready, Almost There, or Needs Preparation — and the candidate’s history.

Design-wise, we used **Strategy** for question setters and evaluators, a **Factory** to pick the best evaluator, a **Decorator** for the fallback, an **Observer-style listener** to keep the engine independent of the console, and a **Bridge** between evaluators and LLM HTTP clients. It is built with **Maven**, packaged as an executable JAR with the shade plugin, and verified with **34 JUnit 5 tests**, all passing.”

1.2 Team

Member	Main responsibility (report §7.1)
RITHIKA S (2104251040779)	Service layer: KeywordBasedEvaluator, AI evaluators (Claude, Gemini, Ollama), FallbackEvaluator, EvaluatorFactory, the three question setters
SHWETHAA JV (2104251050318)	Data layer & testing: questions.json schema, JsonQuestionLoader, InMemoryQuestionBank, the 34 JUnit 5 tests
Supervisor	Dr. R. Kavitha, M.E., Ph.D., Associate Professor, Dept. of CSE

1.3 Problem, driving question and objectives

Problem: Students prepare for core-Java technical rounds with static notes/quizzes — no interactive practice, no instant feedback, no topic-wise diagnosis.

Driving question (from the report): “Can we build a Java console application that conducts a realistic, topic-focused core-Java mock interview, scores answers automatically, and reports placement readiness, without requiring a network connection or a paid service?”

Objectives:

1. Design a question-bank schema and load it from an external JSON file.
2. Model the system with OOP classes and interfaces for question selection and answer evaluation.
3. Implement offline keyword scoring, with an optional AI-backed evaluator behind the same interface.
4. Run a topic-focused mock interview and report placement readiness — no network or paid service needed.
5. Verify the core logic with an automated JUnit 5 test suite.

1.4 Features list

Feature	How it is implemented
Candidate registration	<code>CandidateRegistrationUI.register()</code> reads name, university, roll no. (blank → "Unknown") and creates a <code>Candidate</code> .
Main menu (5 options)	<code>InterviewSimulatorApp.run()</code> — a <code>while</code> loop with a Java 14+ switch expression .
3 question-setting strategies	<code>RandomQuestionSetter</code> , <code>BalancedQuestionSetter</code> , <code>TopicQuestionSetter</code> implementing <code>QuestionSetter</code> .
Offline evaluation	<code>KeywordBasedEvaluator</code> : $\text{score} = \text{matched keywords} \div \text{total keywords} \times 100$.
AI evaluation (optional)	<code>ClaudeAnswerEvaluator</code> , <code>GeminiAnswerEvaluator</code> , <code>OllamaAnswerEvaluator</code> via <code>AbstractLlmEvaluator</code> + <code>LlmClient</code> .
Automatic fallback	<code>FallbackEvaluator</code> wraps AI + keyword evaluator; switches on any <code>RuntimeException</code> .
Per-answer feedback	Score, feedback, hit keywords, missed keywords and the model answer (<code>ConsoleInterviewUI.onAnswerEvaluated</code>).
Readiness report	<code>PlacementReadinessReportGenerator</code> : overall %, per-topic % with ASCII bars, verdict (≥ 75 Ready, ≥ 50 Almost There), history.
Browse question bank	Menu 2: list questions of a chosen topic.
Profile & statistics	Menu 3: <code>Candidate.toString()</code> — interviews completed, best score.
AI question generation	Menu 4: <code>AiQuestionGenerator</code> asks the LLM for 3 new questions as a JSON array and adds them to the bank at runtime.

1.5 Question bank distribution

Topic	OOP	Collections	Strings	Exceptions	JVM	Multithreading	Total
Questions	4	4	3	3	4	5	23

Each JSON entry has five fields: `id`, `topic`, `question`, `keywords` (array) and `modelAnswer`.

```
{
  "id": 1,
  "topic": "OOP",
  "question": "What is the difference between abstract classes and interfaces in Java?",
  "keywords": ["abstract", "interface", "multiple inheritance", "methods", "extend", "implement"],
  "modelAnswer": "An abstract class can hold both abstract and concrete methods ..."
}
```

2. Technology stack & build

Item	Choice	Why / what to say
Language	Java 17 (LTS)	Long-term support release; lets us use switch expressions, text blocks, pattern matching for <code>instanceof</code> , <code>var</code> , <code>String.isBlank()/repeat()</code> , and the built-in <code>java.net.http.HttpClient</code> .
Build tool	Apache Maven	Declarative dependency management (<code>pom.xml</code>), standard directory layout, one-command build/test/package.
JSON library	Gson 2.11.0	Lightweight (single JAR, no transitive deps), simple reflection-based mapping to POJOs, easy custom deserializers.
Testing	JUnit Jupiter 5.10.2	Standard Java unit-testing framework; run by maven-surefire-plugin on <code>mvn test</code> .
Packaging	maven-shade-plugin 3.5.2	Builds a “fat”/uber JAR that bundles Gson and sets <code>Main-Class</code> in the manifest → <code>java -jar interview-simulator-1.0.0.jar</code> .
HTTP	<code>java.net.http.HttpClient</code>	Part of the JDK since Java 11 — no extra library needed for calling AI APIs.
AI providers (optional)	Claude, Gemini, Ollama	Configured purely through environment variables; app works fully offline without them.
Storage	None (in-memory)	Questions in an <code>ArrayList</code> ; candidate stats live for one run only.
IDE / VCS	VS Code / IntelliJ IDEA, Git	Feature branches and commits per milestone; repo github.com/Rithikas07/interview-simulator .
Hardware (report Table 3.2)	Intel i5 or equivalent, 8 GB RAM, Windows 11	Any modest laptop; Maven 3.9.6 used for builds.

2.1 Commands you should know

```
# compile and run all 34 tests
mvn test
# build the executable fat JAR into target/
mvn clean package
# run it
java -jar target/interview-simulator-1.0.0.jar

# optional AI configuration (environment variables)
ANTHROPIC_API_KEY=... ANTHROPIC_MODEL=claude-sonnet-4-5 (default model)
GEMINI_API_KEY=... GEMINI_MODEL=gemini-2.5-flash
OLLAMA_URL=http://localhost:11434 OLLAMA_MODEL=llama3.1
```

2.2 Project structure (Maven standard layout)

```
interview-simulator/
├── pom.xml
└── src/
    ├── main/java/com/interviewsim/
    │   ├── InterviewSimulatorApp.java    ← main(), menu, wiring
    │   ├── core/                        InterviewSession, InterviewListener
    │   ├── model/                       Candidate, Question, Topic, AnswerEvaluation
    │   ├── service/                     QuestionBank, QuestionSetter(+3), AnswerEvaluator(+...),
    │   │                                   LlmClient(+3), EvaluatorFactory, ReportGenerator(+1), ...
    │   ├── ui/                          CandidateRegistrationUI, ConsoleInterviewUI
    │   └── util/                         JsonQuestionLoader
    ├── main/resources/questions.json     ← packaged inside the JAR (classpath)
    └── test/java/com/interviewsim/      ← 11 JUnit 5 test classes
```

2.3 Maven build lifecycle (frequently asked)

validate → compile → test → package → verify → install → deploy. Running a phase runs every phase before it, so `mvn package` also compiles and runs tests. Plugins bind goals to phases: surefire's `test` goal runs in `test`; shade's `shade` goal is bound to `package` in our `pom.xml`.

3. System architecture

A layered architecture with a strict top-to-bottom dependency direction. Every layer talks to the one below through **interfaces**, so implementations can be swapped without touching callers.




Figure 1 — Layered architecture (packages, key classes and external services)

3.1 Layer responsibilities

Layer / package	Responsibility	Classes
Presentation (ui + app)	All console input/output; menu loop; wiring of objects (composition root).	InterviewSimulatorApp, CandidateRegistrationUI, ConsoleInterviewUI
Core (core)	The interview engine — UI-agnostic. Pulls questions, asks via listener, evaluates, averages, records result.	InterviewSession, InterviewListener
Service (service)	Business logic behind interfaces: question supply, selection strategies, evaluation, AI connectivity, reporting.	QuestionBank, QuestionSetter, AnswerEvaluator, LlmClient, ReportGenerator + implementations
Model (model)	Plain data / domain objects with encapsulated state.	Candidate, Question, Topic, AnswerEvaluation
Util / data	Loading the JSON question bank from the classpath.	JsonQuestionLoader, questions.json

3.2 The pipeline (as shown on the poster)

Candidate Registration → Question-Setting Strategy Selection → Question Bank (JSON via Gson) → Answer Evaluation (AI + Fallback / Keyword) → Placement-Readiness Report

3.3 Why this architecture? (likely question)

- **Separation of concerns:** the engine never calls `System.out`; the UI never computes scores.
- **Testability:** every service can be unit-tested without a console or network (e.g. `FallbackEvaluatorTest` uses a fake AI that throws).
- **Extensibility:** a new topic is a new enum constant + JSON entries; a new AI provider is one `LlmClient` + one evaluator subclass; a GUI is a new `InterviewListener`.
- **Resilience:** the system degrades gracefully to offline scoring.

3.4 Object wiring at startup (composition root)

```
public InterviewSimulatorApp() {
    this.candidate = new CandidateRegistrationUI(scanner).register();
    this.evaluator = EvaluatorFactory.selectBestEvaluator(); // Factory picks the chain
}
// fields: Scanner scanner = new Scanner(System.in);
//         QuestionBank questionBank = new InMemoryQuestionBank(); ← Loads questions.json

// on menu 1:
InterviewSession session = new InterviewSession(questionBank, evaluator, candidate, setter);
List<AnswerEvaluation> evaluations = session.run(new ConsoleInterviewUI(scanner));
System.out.println(new PlacementReadinessReportGenerator(candidate).generateReport(candidate, evaluations));
```

Dependencies are passed into `InterviewSession` through its constructor — this is **constructor-based dependency injection** (done by hand, without a framework like Spring).

4. Class diagrams

4.1 Interview engine & question subsystem




Figure 2 — InterviewSession depends only on interfaces; question setters delegate to RandomQuestionSetter

4.2 Evaluation & AI subsystem




Figure 3 — Evaluator hierarchy: Strategy + Decorator (FallbackEvaluator) + Bridge (AbstractLlmEvaluator ↔ LlmClient)

4.3 Model classes

Class	Fields	Notes for viva
Candidate	<code>final</code> name, university, rollNumber; bestScorePercentage, interviewsCompleted	Mutable stats, immutable identity. Sentinel <code>NO_SCORE_YET = -1</code> . <code>recordInterviewResult()</code> updates both stats. Overrides <code>toString()</code> .
Question	<code>final</code> id, topic, text, expectedKeywords, modelAnswer	<code>@SerializedName</code> maps JSON keys (e.g. "question" → text). Private no-arg constructor for Gson. Returns <code>Collections.unmodifiableList</code> .
Topic (enum)	displayName	Enum with constructor and field: OOP, COLLECTIONS, STRINGS, EXCEPTIONS, JVM, MULTITHREADING.
AnswerEvaluation	question, candidateAnswer, score, matchedKeywords, missedKeywords, feedback (all <code>final</code>)	Immutable value object ; defensive copies via <code>List.copyOf</code> . No setters.

4.4 Interfaces at a glance

Interface	Methods	Implemented by
QuestionBank	getAllQuestions, getRandomQuestions, getQuestionsByTopic, addQuestions	InMemoryQuestionBank
QuestionSetter	selectQuestions(bank,count), selectQuestions(bank,topic,count), getStrategyName	Random-, Balanced-, TopicQuestionSetter
AnswerEvaluator	evaluate(Question, String)	KeywordBasedEvaluator, FallbackEvaluator, (via AiAnswerEvaluator) AI evaluators
AiAnswerEvaluator <i>extends</i> AnswerEvaluator	+ isAvailable()	AbstractLlmEvaluator → Claude/Gemini/Ollama
LlmClient	isAvailable, chat(system,user), getProviderName	ClaudeLlmClient, GeminiLlmClient, OllamaLlmClient
ReportGenerator	generateReport(Candidate, List<AnswerEvaluation>)	PlacementReadinessReportGenerator
InterviewListener	onSessionStart, onAskQuestion, onAnswerEvaluated, onSessionEnd	ConsoleInterviewUI

5. Flow diagrams

5.1 Application flow (main menu)




Figure 4 — Overall application flowchart (InterviewSimulatorApp.run)

5.2 Interview session — sequence diagram




Figure 5 — Sequence of calls for one interview (InterviewSession.run)

5.3 Evaluator selection at startup




Figure 6 — EvaluatorFactory.selectBestEvaluator()

5.4 Runtime fallback inside FallbackEvaluator




Figure 7 — `FallbackEvaluator.evaluate()`

5.5 AI evaluation pipeline




Figure 8 — `AbstractLlmEvaluator.evaluate()` → `parseEvaluation()`

5.6 Question bank loading




Figure 9 — `JsonQuestionLoader.load()` at startup

6. Design patterns used


A **design pattern** is a proven, reusable solution to a commonly occurring design problem — a template for how to arrange classes and objects, not finished code. The classic catalogue is the “Gang of Four” (GoF) book *Design Patterns* (Gamma, Helm, Johnson, Vlissides, 1994), cited on our poster. GoF groups patterns into three families:

Family	Concerned with	In our project
Creational	How objects are created	Factory (EvaluatorFactory), Builder (HttpClient, HttpRequest, GsonBuilder — used from libraries)
Structural	How classes/objects are composed into larger structures	Decorator (FallbackEvaluator), Bridge (AbstractLlmEvaluator ↔ LlmClient), Adapter (each LlmClient)
Behavioural	How objects communicate and share responsibility	Strategy (QuestionSetter, AnswerEvaluator), Observer/Listener (InterviewListener), Template Method (AbstractLlmEvaluator), Chain-of-responsibility-like fallback
Idioms	Smaller Java-level practices	Dependency Injection, Immutable Value Object, Utility class, Delegation (composition over inheritance)

6.1 Strategy pattern Behavioural

What the pattern is: Define a *family of algorithms*, put each in its own class behind a common interface, and make them *interchangeable at runtime*. The client (“context”) holds a reference to the interface and does not know which concrete algorithm it is using. It replaces long `if/else` or `switch` blocks that choose behaviour.

Participants: *Strategy* (interface) · *ConcreteStrategy* (implementations) · *Context* (uses a strategy).



Where in our project — twice:

- Question selection:** `QuestionSetter` (strategy) → `RandomQuestionSetter`, `BalancedQuestionSetter`, `TopicQuestionSetter`. The user picks one at runtime from the menu.
- Answer evaluation:** `AnswerEvaluator` → `KeywordBasedEvaluator`, AI evaluators, `FallbackEvaluator`.
- Context:** `InterviewSession` — it only calls `questionSetter.selectQuestions(...)` and `evaluator.evaluate(...)`.

```
QuestionSetter setter = switch (choice) {           // chooseQuestionSetter()
    case 1 -> new RandomQuestionSetter();
    case 2 -> new BalancedQuestionSetter();
    case 3 -> new TopicQuestionSetter(topic);
    default -> null;
};
// InterviewSession never knows which one:
List<Question> questions = questionSetter.selectQuestions(questionBank, questionCount);
```

Benefit: Adding a “Difficulty-based” setter needs one new class — `InterviewSession` is untouched (Open/Closed principle). Each strategy is unit-tested separately (`RandomQuestionSetterTest`, `BalancedQuestionSetterTest`,

6.2 Factory pattern (Simple / Static Factory) Creational

What the pattern is: Move object-creation logic out of the client into a dedicated *factory* that decides *which* concrete class to instantiate and returns it through an interface type. The client asks for “an evaluator” and does not use `new` on concrete classes. (GoF’s *Factory Method* uses subclass overriding; what we have is the simpler **Simple Factory / static factory method** variant — be ready to state this distinction.)


Where: `EvaluatorFactory` — a `final` class with a `private` constructor and static methods `getAiEvaluators()`, `getStaticEvaluator()`, `selectBestEvaluator()`, `describeSelectedEvaluator()`.

```
public static AnswerEvaluator selectBestEvaluator() {
    for (AiAnswerEvaluator ai : getAiEvaluators()) { // Claude → Gemini → OLLama
        if (ai.isAvailable()) {
            return new FallbackEvaluator(ai, getStaticEvaluator());
        }
    }
    return getStaticEvaluator(); // KeywordBasedEvaluator
}
```

Benefit: The app’s constructor is one line (`EvaluatorFactory.selectBestEvaluator()`). Priority order and wiring of the decorator live in one place. Return type is the interface, so the caller is decoupled from concrete classes.

6.3 Decorator pattern Structural

What the pattern is: Attach extra responsibilities to an object *dynamically* by wrapping it in another object that **implements the same interface** and forwards calls to it, adding behaviour before/after/around the call. It is a flexible alternative to subclassing. Classic example: `new BufferedReader(new InputStreamReader(in))` in `java.io`.



Where: `FallbackEvaluator` implements `AnswerEvaluator` and wraps an `AiAnswerEvaluator` (primary). It adds **resilience**: if the AI throws, it disables AI and delegates to the keyword evaluator.

```
public AnswerEvaluation evaluate(Question question, String candidateAnswer) {
    if (aiActive) {
        try {
            return primary.evaluate(question, candidateAnswer); // forward
        } catch (RuntimeException e) {
            aiActive = false; // added behaviour
            System.err.println("[INFO] AI evaluator unavailable ... Falling back ...");
        }
    }
    return fallback.evaluate(question, candidateAnswer);
}
```

Viva nuance: Strictly, a pure decorator wraps *one* component; ours wraps two and chooses between them, so it is also close to a **Proxy** (controls access to the AI) and a two-link **Chain of Responsibility** (try handler 1, else handler 2). Saying “Decorator, with a chain-of-responsibility flavour” shows depth.

Benefit: InterviewSession and the UI are completely unaware of failures; the interview never crashes because the internet dropped.

6.4 Observer / Listener (callback) pattern Behavioural

What the pattern is: A *subject* notifies one or more *observers* about events through a common interface, without knowing their concrete classes. In Java it often appears as a **Listener** interface (Swing's `ActionListener`, for example).




Where: InterviewListener (in core) with 4 callbacks: `onSessionStart`, `onAskQuestion` (returns the answer), `onAnswerEvaluated`, `onSessionEnd`. `ConsoleInterviewUI` implements it; `InterviewSession.run(listener)` calls it.

Benefit: The engine contains zero `System.out` calls — it can be driven by a console, a Swing GUI, a web front-end or a test stub. (Note: `onAskQuestion` returns a value, so it is a callback interface — “observer-style”, as the Javadoc says — rather than a pure one-way observer.)

6.5 Bridge pattern Structural

What the pattern is: Split a large class into two independent hierarchies — an **abstraction** and an **implementation** — connected by composition (“has-a”), so both can vary independently. Without it you would need one class for every combination (N evaluators × M providers).



Where: `AbstractLlmEvaluator` (abstraction: build prompt, parse verdict, compute keywords) **has-a** `LlmClient` (implementation: HTTP details of Claude / Gemini / Ollama). The same `LlmClient`s are reused by `AiQuestionGenerator` — proof that the split is useful.

```
public class ClaudeAnswerEvaluator extends AbstractLlmEvaluator {
    public ClaudeAnswerEvaluator() { super(new ClaudeLlmClient()); } // that's all!
}
```

6.6 Template Method pattern Behavioural

What the pattern is: An abstract base class defines the *skeleton of an algorithm* in a method, and lets subclasses override certain *steps* (“hooks”) without changing the overall structure.

Where: `AbstractLlmEvaluator.evaluate()` is the fixed skeleton: check availability → `chat(getSystemPrompt(), buildUserPrompt())` → `parseEvaluation()` → `extractJson()`. The protected methods `getSystemPrompt()`,

`parseEvaluation()` and `extractJson()` are hooks a subclass may override (e.g. a provider-specific prompt). Our three subclasses currently reuse all defaults.

6.7 Adapter pattern Structural

What the pattern is: Convert the interface of an existing class/service into the interface clients expect, so incompatible things can work together (like a travel plug adapter).

Where: Each `LlmClient` adapts a *different* remote API to the same simple `chat(systemPrompt, userPrompt) → String` method:

Client	Endpoint	Request shape it adapts to	Reply path it unwraps
ClaudeLlmClient	api.anthropic.com/v1/messages	model, max_tokens, system, messages[] + headers x-api-key, anthropic-version	content[0].text
GeminiLlmClient	.../models/{model}:generateContent? key=...	contents[].parts[], systemInstruction, generationConfig	candidates[0].content.parts[0].text
OllamaLlmClient	localhost:11434/api/chat	model, stream:false, messages[], options	message.content

6.8 Builder pattern (used, from libraries) Creational

What the pattern is: Construct a complex object step by step through a fluent builder, then call `build()`. Avoids “telescoping constructors” with many parameters and produces immutable objects.

Where we use it: `HttpClient.newBuilder().connectTimeout(...).build()`, `HttpRequest.newBuilder().uri(...).timeout(...).header(...).POST(...).build()`, and `new GsonBuilder().registerTypeAdapter(...).create()`. Also `StringBuilder` in the report generator (builder-like for strings). We use the pattern; we did not write our own builder class.

6.9 Supporting idioms & principles Idioms

Idiom	What it is	Where
Dependency Injection (constructor)	An object receives its collaborators from outside instead of creating them, which makes it testable and configurable.	<code>InterviewSession(bank, evaluator, candidate, setter);</code> <code>FallbackEvaluator(primary, fallback);</code> <code>RandomQuestionSetter(Random)</code> for seeded tests.
Immutable Value Object	Object whose state never changes after construction — thread-safe and safe to share.	<code>AnswerEvaluation</code> (all <code>final</code> , <code>List.copyOf</code> , no setters).
Delegation / composition over inheritance	Reuse behaviour by holding and calling another object rather than extending it.	<code>BalancedQuestionSetter</code> and <code>TopicQuestionSetter</code> hold a <code>RandomQuestionSetter</code> delegate.

Utility / static-factory class	<code>final</code> class with <code>private</code> constructor — cannot be instantiated or subclassed.	<code>EvaluatorFactory</code> , <code>JsonQuestionLoader</code> is <code>final</code> .
Null-object-like defaults	Return safe defaults instead of <code>null</code> .	Loader returns <code>List.of()</code> on error; registration returns “Unknown” for blanks.
Custom exception wrapping	Translate low-level exceptions to one domain exception (exception chaining).	<code>AiEvaluationException</code> wraps <code>IOException</code> , <code>InterruptedException</code> , parse errors.

6.10 Pattern quick-recall table

Pattern	Class(es)	One-line justification
Strategy	<code>QuestionSetter</code> ×3, <code>AnswerEvaluator</code> ×5	Swap selection/scoring algorithms at runtime without changing <code>InterviewSession</code> .
Factory	<code>EvaluatorFactory</code>	Central place to decide which evaluator chain to build.
Decorator	<code>FallbackEvaluator</code>	Adds fault-tolerance around the AI evaluator through the same interface.
Observer / Listener	<code>InterviewListener</code> , <code>ConsoleInterviewUI</code>	Engine notifies UI; UI can be replaced.
Bridge	<code>AbstractLlmEvaluator</code> ↔ <code>LlmClient</code>	Evaluation logic and provider HTTP logic vary independently.
Template Method	<code>AbstractLlmEvaluator.evaluate()</code>	Fixed algorithm skeleton with overridable hooks.
Adapter	<code>ClaudeLlmClient</code> , <code>GeminiLlmClient</code> , <code>OllamaLlmClient</code>	Three different REST APIs behind one <code>chat()</code> .
Builder	<code>HttpClient</code> , <code>HttpRequest</code> , <code>GsonBuilder</code>	Fluent construction of complex library objects.

Common follow-up: “Why not Singleton for `EvaluatorFactory`?” — It has no state to share; static methods are enough. A Singleton is for exactly-one *instance* with state (e.g. a connection pool). We deliberately avoided Singleton because global state hurts testability.

7. The four OOP pillars in our project

Pillar	Definition	Concrete example in our code
Encapsulation	Bundle data with the methods that operate on it; hide internal state behind access control.	Candidate's fields are <code>private</code> ; best score and count can only change through <code>recordInterviewResult()</code> , which enforces the rule "keep the maximum". <code>Question.getExpectedKeywords()</code> returns an unmodifiable view so callers can't mutate internal lists. <code>FallbackEvaluator.aiActive</code> is private with a read-only getter.
Abstraction	Expose <i>what</i> an object does, hide <i>how</i> .	Interfaces <code>QuestionBank</code> , <code>QuestionSetter</code> , <code>AnswerEvaluator</code> , <code>LlmClient</code> , <code>ReportGenerator</code> , <code>InterviewListener</code> ; the abstract class <code>AbstractLlmEvaluator</code> . <code>InterviewSession</code> only sees abstractions.
Inheritance	A class acquires members of a parent (<code>extends</code>) or a type contract (<code>implements</code>).	<code>ClaudeAnswerEvaluator</code> extends <code>AbstractLlmEvaluator</code> ; <code>AiAnswerEvaluator</code> extends <code>AnswerEvaluator</code> (interface inheritance); <code>AiEvaluationException</code> extends <code>RuntimeException</code> .
Polymorphism	One interface, many forms. <i>Runtime</i> (overriding, dynamic dispatch) and <i>compile-time</i> (overloading).	Runtime: <code>evaluator.evaluate(q, a)</code> runs <code>Keyword</code> / <code>Fallback</code> / <code>Claude</code> code depending on the actual object. Compile-time: overloaded <code>selectQuestions(bank, count)</code> vs <code>selectQuestions(bank, topic, count)</code> ; overloaded constructors in <code>InterviewSession</code> , <code>RandomQuestionSetter</code> , <code>ClaudeLlmClient</code> .

7.1 Other OOP concepts visible in the code

- **Constructor chaining:** `this(...)` in `InterviewSession`, `ClaudeLlmClient()`, `TopicQuestionSetter()`; `super(...)` in evaluator subclasses and the exception class.
- **Access modifiers:** `private` fields everywhere, `protected` hooks/fields in `AbstractLlmEvaluator` (`gson`, `parseEvaluation`), `public` API methods, `private` constructor in `EvaluatorFactory`.
- **Abstract class vs interface:** both used — interface for contracts, abstract class for shared code (prompt + parsing).
- **Association / aggregation / composition:** `InterviewSession` *has* a bank, evaluator, setter (aggregation — they live beyond the session); `AnswerEvaluation` *has* a `Question`.
- **@Override** annotation on every overriding method → compiler checks signatures.
- **static and final:** constants such as `DEFAULT_QUESTION_COUNT = 5`, `READY_THRESHOLD = 75`; nested `static class` stubs in tests.
- **Enum as a class:** `Topic` has a private constructor, a field and a getter.
- **instanceof and runtime type checks:** `findClientFor()` in the app; pattern matching in `EvaluatorFactory`.

8. SOLID principles in the project

Principle	Meaning	Evidence in our code
S — Single Responsibility	A class should have one reason to change.	<code>JsonQuestionLoader</code> only loads JSON; <code>InMemoryQuestionBank</code> only stores/queries; <code>KeywordBasedEvaluator</code> only scores; <code>PlacementReadinessReportGenerator</code> only formats the report; <code>ConsoleInterviewUI</code> only does I/O.
O — Open/Closed	Open for extension, closed for modification.	New question setter / evaluator / LLM provider = new class; <code>InterviewSession</code> is not modified. New topic = enum constant + JSON.
L — Liskov Substitution	Subtypes must be usable wherever the parent type is expected.	Any <code>AnswerEvaluator</code> (Keyword, Fallback, Claude...) can be passed to <code>InterviewSession</code> ; tests substitute fake <code>AiAnswerEvaluator</code> s.
I — Interface Segregation	Clients shouldn't depend on methods they don't use.	<code>isAvailable()</code> is in <code>AiAnswerEvaluator</code> , not in <code>AnswerEvaluator</code> — the keyword evaluator isn't forced to implement an availability check it doesn't need.
D — Dependency Inversion	Depend on abstractions, not concretions.	<code>InterviewSession</code> fields are interface-typed (<code>QuestionBank</code> , <code>AnswerEvaluator</code> , <code>QuestionSetter</code>); <code>AbstractLlmEvaluator</code> depends on <code>LlmClient</code> , not on HTTP classes.

Be honest if asked where SOLID is violated: `InterviewSimulatorApp.findClientFor()` uses an `instanceof` chain to map an evaluator to its client — adding a provider means editing this method (an OCP violation). A cleaner fix: add `LlmClient getClient()` to `AiAnswerEvaluator`, or let the factory return the client too.

9. Class-by-class walkthrough

Class	Type	What it does (say this in the viva)
com.interviewsim		
InterviewSimulatorApp	class (main)	Entry point. Creates Scanner, question bank, registers candidate, asks factory for evaluator; runs the 5-option menu; builds InterviewSession for each interview; prints report; handles AI question generation; <code>readInt()</code> returns -1 on bad input.
core		
InterviewSession	class	Engine. <code>run(listener)</code> : select questions → start callback → for each: ask, evaluate, notify → average score → <code>candidate.recordInterviewResult()</code> → end callback → return evaluations. Default 5 questions.
InterviewListener	interface	Callback contract between engine and UI (4 methods).
model		
Candidate	class	Profile + stats; <code>recordInterviewResult()</code> increments count and keeps max score.
Question	class	Immutable question; Gson-mapped with <code>@SerializedName</code> .
Topic	enum	6 topics with display names.
AnswerEvaluation	class	Immutable result: score, matched/missed keywords, feedback.
service — questions		
QuestionBank / InMemoryQuestionBank	interface / class	Stores questions in an <code>ArrayList</code> loaded via <code>JsonQuestionLoader</code> ; returns copies of lists (defensive); filters by topic; can append AI-generated questions.
QuestionSetter	interface	Strategy for choosing questions; has overloaded methods + <code>getStrategyName()</code> .
RandomQuestionSetter	class	Copy → <code>Collections.shuffle(copy, random)</code> → first <i>n</i> . Injectable <code>Random</code> for repeatable tests.
BalancedQuestionSetter	class	Round-robin over shuffled per-topic pools (see §10.2).
TopicQuestionSetter	class	Bound to one topic (default OOP); delegates to Random with that topic.
service — evaluation & AI		
AnswerEvaluator	interface	<code>evaluate(Question, String) → AnswerEvaluation</code> .
KeywordBasedEvaluator	class	Case-insensitive <code>contains</code> check per keyword; score = <code>round(matched × 100 / total)</code> ; feedback by thresholds 75/40.
AiAnswerEvaluator	interface	Extends AnswerEvaluator with <code>isAvailable()</code> .
AbstractLlmEvaluator	abstract class	System prompt with scoring rubric (text block), builds user prompt, calls <code>LlmClient</code> , extracts/parses JSON verdict, clamps score, adds keyword lists, prefixes "[AI EVALUATOR]".
Claude/Gemini/OllamaAnswerEvaluator	classes	Thin subclasses that pass the right <code>LlmClient</code> to <code>super(...)</code> .
LlmClient + 3 clients	interface / classes	HTTP adapters. Claude & Gemini "available" if API key env var set; Ollama available if <code>GET /api/tags</code> returns 200 within 3 s.
FallbackEvaluator	class	Decorator: try AI, on exception switch permanently to keyword evaluator.

EvaluatorFactory	final class	Static factory with priority Claude → Gemini → Ollama → keyword.
AIException	class	Unchecked custom exception (extends RuntimeException) with cause-chaining constructor.
AIQuestionGenerator	class	Prompts the LLM for a JSON array of questions; parses with <code>TypeToken<List<JsonObject>></code> ; skips malformed entries; assigns ids from 1000.
service — reporting		
ReportGenerator / PlacementReadinessReportGenerator	interface / class	Builds report with <code>StringBuilder</code> : header, overall, per-topic average via <code>Map<Topic,int[]> + computeIfAbsent</code> , ASCII bar (<code>"#".repeat(n)</code>), verdict, history.
ui & util		
CandidateRegistrationUI	class	Prompts 3 fields; blank → "Unknown".
ConsoleInterviewUI	class	Implements InterviewListener with formatted console output.
JsonQuestionLoader	final class	Loads <code>/questions.json</code> from classpath with Gson + custom Topic deserializer (lambda); never throws — returns empty list on error.

10. Key algorithms explained

10.1 Keyword-based scoring

```
normalized = answer == null ? "" : answer.toLowerCase();
for (String keyword : question.getExpectedKeywords())
    (normalized.contains(keyword.toLowerCase()) ? matched : missed).add(keyword);
score = total == 0 ? 0 : Math.round(matched.size() * 100.0 / total);
```

Worked example: Q “Four pillars of OOP”, keywords = [encapsulation, inheritance, polymorphism, abstraction]. Answer: “Encapsulation hides data and inheritance reuses code.” → matched 2/4 → **50** → “Partially correct — covered 2/4 key points.” (≥75 Excellent, ≥40 Partial, else Weak; blank → “No answer given”).

Complexity: $O(k \times n)$ where k = number of keywords and n = answer length (each `contains` is a linear scan).

10.2 Balanced (round-robin) question selection




Figure 10 — `BalancedQuestionSetter` picks one question per topic per round

- For each `Topic.values()`, copy its questions into a pool and shuffle it; skip empty pools.
- Keep an `int[]` cursor — one index per pool.
- Loop rounds: for each pool, if `cursor[i] < pool.size()` take `pool.get(cursor[i]++)`; stop when `result.size() == count`.
- If a full round adds nothing (`addedThisRound == false`), all pools are exhausted → `break` (prevents an infinite loop when count > bank size).

With the default 5 questions you get one question each from 5 different topics (OOP, Collections, Strings, Exceptions, JVM — in enum order). Complexity $O(N)$ for N questions (plus shuffling).

10.3 Session score, topic average and verdict

- Session score** = `Math.round(sum of scores / number of questions)` (each score 0–100).
- Topic average:** `Map<Topic, int[]>` where `int[0]` = sum and `int[1]` = count, filled using `computeIfAbsent(topic, t -> new int[2])`.
- Bar:** `filled = round(pct/10) → "#".repeat(filled) + "-".repeat(10 - filled)`, e.g. 73% → `#####---`.
- Verdict:** ≥ 75 → READY; ≥ 50 → ALMOST THERE; else NEEDS PREPARATION. Poster demo: 73% → ALMOST THERE.
- Best score:** `Candidate.recordInterviewResult()` keeps the maximum; `-1` means “no interview yet” (printed “N/A”).

10.4 AI verdict parsing (robustness)

1. System prompt (a Java text block) tells the model to reply **only** with `{"score": 0-100, "feedback": "..."}` and gives a rubric (90–100 comprehensive ... 0 empty).
2. LLMs sometimes add prose or code fences, so `extractJson()` takes the substring from the first `{` to the last `}`.
3. Gson parses it to a `JsonObject`; score is **clamped** with `Math.max(0, Math.min(100, raw))`.
4. Matched/missed keywords are still computed locally so the UI shows the same details as offline mode.
5. Any parse failure → `AIException` → `FallbackEvaluator` takes over.

10.5 AI question generation

Prompt asks for a JSON **array** of objects `{topic, question, keywords[], modelAnswer}`. The code extracts `[...]`, parses with `new TypeToken<List<JsonObject>>().getType()`, maps topic text with `Topic.valueOf(s.trim().toUpperCase().replace(' ', '_'))`, and **skips malformed entries** instead of failing the whole batch. New questions get ids from 1000 and are added with `questionBank.addQuestions()`.

11. Project viva questions & answers

11.1 General / motivation

Q1. What is your project? Explain it briefly.

See the 2-minute pitch in §1.1. Short version: a Java 17 console app that conducts a 5-question core-Java mock interview from a JSON question bank, scores answers using AI (Claude/Gemini/Ollama) with automatic fallback to offline keyword scoring, and prints a topic-wise placement-readiness report.

Q2. What problem does it solve? Who is the user?

Engineering students preparing for campus placements. Static notes give no feedback; our tool gives instant per-answer scoring, the model answer, missed keywords and a topic-wise diagnosis, offline and free.

Q3. Why a console application and not a GUI or web app?

Scope and time: two members, one PBL cycle. Console keeps focus on core-Java concepts (OOP, collections, exceptions) — the subject being examined. The architecture is UI-agnostic (InterviewListener), so a Swing/JavaFX/web front-end can be added later without changing the engine. We considered Spring Boot + Hibernate + REST but rejected it for setup complexity (report §2.1.2).

Q4. How is your project different from just asking ChatGPT?

(1) Works fully offline with keyword scoring; (2) structured, curated questions with model answers and expected keywords; (3) consistent scoring rubric; (4) topic-wise readiness report and best-score tracking; (5) choice of question strategies. The AI is an optional plug-in, not a dependency.

Q5. Why did you choose Java 17?

It is an LTS release, and it gives us switch expressions, text blocks, `instanceof` pattern matching, `var`, `List.of/copyOf`, `String.isBlank/repeat` and the standard `HttpClient` — all used in our code.

Q6. Why Gson and not Jackson or org.json?

Gson is a single small JAR with no transitive dependencies, maps JSON to POJOs with minimal configuration, and supports custom deserializers in one lambda. Jackson is more powerful/faster but heavier than needed for a 23-question file; org.json would require manual mapping.

11.2 Architecture & design

Q7. Explain the architecture of your project.

Five layers (Fig. 1): presentation (App + UI), core engine (InterviewSession + InterviewListener), service (question supply, evaluation, AI connectivity, reporting — all behind interfaces), model (Candidate, Question, Topic, AnswerEvaluation) and util/data (JsonQuestionLoader + questions.json). Dependencies flow downward only.

Q8. Which design patterns have you used and why?

Strategy (question setters, evaluators), Factory (EvaluatorFactory), Decorator (FallbackEvaluator), Observer/Listener (InterviewListener), Bridge (AbstractLlmEvaluator ↔ LlmClient), Template Method (AbstractLlmEvaluator.evaluate), Adapter (each LlmClient), and we use Builders from the JDK/Gson. See chapter 6 for justification of each.

Q9. Why is InterviewListener an interface? Why not just print inside InterviewSession?

To decouple the engine from the UI. InterviewSession can then be tested without a console and reused with a GUI or web front-end. It follows the Dependency Inversion principle and the observer/callback pattern.

Q10. Why do you have both `AnswerEvaluator` and `AiAnswerEvaluator`?

Interface Segregation. Only AI evaluators need `isAvailable()`. The keyword evaluator is always available, so it implements only `AnswerEvaluator`. The factory and `FallbackEvaluator` need the extra method, so they depend on `AiAnswerEvaluator`.

Q11. Why an abstract class `AbstractLlmEvaluator` instead of putting code in each provider evaluator?

The prompt, JSON extraction, score clamping and keyword computation are identical for all providers — the abstract class removes duplication (DRY). Only the HTTP transport differs, which is delegated to the `LlmClient`. As a result each provider evaluator is just a constructor.

Q12. Why is `LlmClient` separate from the evaluator?

Bridge pattern — separates “what we ask the model” from “how we talk to the provider”. It also lets `AiQuestionGenerator` reuse the same clients for question generation.

Q13. Why does `BalancedQuestionSetter` hold a `RandomQuestionSetter` instead of extending it?

Composition over inheritance. `Balanced` “is not a” random setter — it only reuses one behaviour (topic-scoped random selection). Delegation avoids a fragile inheritance hierarchy and lets each class change independently.

Q14. Why do `QuestionSetter` methods exist in two overloaded forms?

`selectQuestions(bank, count)` is the strategy’s main behaviour; `selectQuestions(bank, topic, count)` is a topic-scoped version used for delegation and topic focus. It demonstrates compile-time polymorphism (method overloading).

Q15. What is dependency injection? Did you use a framework?

Providing an object’s dependencies from outside rather than creating them inside. We do it manually through constructors (e.g. `InterviewSession`, `FallbackEvaluator`, `RandomQuestionSetter(Random)`). No framework (Spring/Guice) — the app class acts as the composition root.

11.3 Evaluation & AI

Q16. How does the keyword-based evaluator calculate the score?

Lower-case the answer, check each expected keyword with `contains`, `score = round(matched × 100 / total)`. Feedback: ≥ 75 “Excellent”, ≥ 40 “Partially correct”, else “Weak”; blank answer gets a dedicated message. Null answers are treated as empty (tested in `nullAnswerIsTreatedAsNoMatches`).

Q17. What are the limitations of keyword scoring?

- Correct answers with different wording/synonyms score low (poster: Exceptions scored 0% on a try-with-resources question).
- Substring false positives: “final” is found inside “finally”; “extend” inside “extends”.
- Keyword stuffing (just listing words) gets full marks; no understanding of meaning or negation (“is *not* thread-safe”).
- Spelling mistakes get no credit.

Mitigations: AI evaluator; word-boundary regex; synonym lists; stemming; semantic/embedding similarity (listed under future enhancements).

Q18. How does the AI evaluator work?

It sends a system prompt (strict interviewer, rubric, “respond only with JSON”) and a user prompt containing the question, expected keywords, model answer and candidate answer. It extracts the JSON from the reply, parses score and feedback with `Gson`, clamps the score to 0–100 and returns an `AnswerEvaluation` (Fig. 8).

Q19. What happens if the internet goes down in the middle of an interview?

The AI client throws `AiEvaluationException` (network error). `FallbackEvaluator` catches it, prints an [INFO] message, sets `aiActive = false` and evaluates that same answer with the keyword evaluator. The interview continues; remaining answers are scored offline.

Q20. Why does AI stay disabled after one failure instead of retrying?

Each failed call may block for the timeout (10 s connect, up to 60–180 s request). Retrying every question would make the interview very slow. Trade-off: a single transient error disables AI for the run. Improvement: retry with exponential back-off or a circuit breaker that re-tries after N questions.

Q21. How does the app decide which AI to use?

`EvaluatorFactory.selectBestEvaluator()` checks providers in priority order Claude → Gemini → Ollama and wraps the first available one in a `FallbackEvaluator`. Claude/Gemini are “available” if their API-key environment variables are non-blank; Ollama if `GET http://localhost:11434/api/tags` returns HTTP 200 within 3 seconds. If none, it uses `KeywordBasedEvaluator` alone.

Q22. How are API keys stored? Is it secure?

Read from environment variables via `System.getenv()` — never hard-coded or committed to Git. Claude’s key goes in the `x-api-key` HTTPS header. Gemini’s key is in the URL query string, which is less ideal because URLs can be logged; the header `x-goog-api-key` would be better.

Q23. Why ask the LLM for JSON? What if it doesn’t comply?

Structured output is machine-parsable. If the model adds extra text, `extractJson()` takes first “{” to last “}”. If it’s still invalid, or the score is out of range, we throw/clamp; any exception triggers fallback.

Q24. What does Ollama offer compared to Claude/Gemini?

Ollama runs open models (default `llama3.1`) locally — free, private, no API key, works without internet once the model is downloaded; but slower on a laptop (hence the 180 s timeout).

Q25. How does AI question generation work, and how do you protect against bad output?

Menu 4 asks the first available LLM for 3 questions as a JSON array. We extract the array, parse each object, map the topic name to the enum, and wrap each entry in try/catch so one malformed entry is skipped with a warning while the others are kept. Then `bank.addQuestions()` appends them for the rest of the run.

Q26. Why is `AiEvaluationException` unchecked (extends `RuntimeException`)?

The `AnswerEvaluator.evaluate()` interface does not declare checked exceptions, so implementations can’t throw checked ones. Unchecked exceptions keep the interface clean; the single place that cares (`FallbackEvaluator`) catches it. We keep the original cause via the `(message, cause)` constructor — exception chaining.

Q27. Why call `Thread.currentThread().interrupt()` when catching `InterruptedException`?

Catching `InterruptedException` clears the thread’s interrupted flag. Re-setting it preserves the interruption status so higher-level code can still see that the thread was asked to stop. Swallowing it is a well-known bad practice.

11.4 Data, model & JSON

Q28. Where are the questions stored and how are they loaded?

In `src/main/resources/questions.json`, which Maven packages inside the JAR. `JsonQuestionLoader` uses `getClass().getResourceAsStream("/questions.json")`, wraps it in a UTF-8 `InputStreamReader` inside try-with-resources, and parses with `Gson` into `List<Question>`. `InMemoryQuestionBank`’s constructor calls it.

Q29. Why getResourceAsStream and not new File("questions.json")?

A file inside a JAR is not a file on disk; `File` paths break once packaged. Classpath resources work both in the IDE and from the fat JAR.

Q30. What is TypeToken and why is it needed?

Because of **type erasure**, `List<Question>.class` doesn't exist at runtime — only `List`. `new TypeToken<List<Question>>(){}` creates an anonymous subclass whose generic superclass information is kept in the class file, so Gson can recover the full type `List<Question>`.

Q31. What does @SerializedName do?

Maps a JSON key to a differently named Java field — e.g. JSON `"question"` → field `text`, `"keywords"` → `expectedKeywords`. Field names can stay meaningful in Java while the JSON stays simple.

Q32. How is the topic string converted into the Topic enum?

A custom `JsonDeserializer<Topic>` registered on a `GsonBuilder`, written as a lambda: `(json, type, ctx) -> Topic.valueOf(json.getAsString().toUpperCase())`. This makes the JSON case-insensitive ("oop" works).

Q33. Why does Question have a private no-arg constructor?

Gson instantiates objects reflectively. If a no-arg constructor exists (even private) Gson uses it, then sets the fields — including `final` ones — via reflection. Without it, Gson falls back to `Unsafe.allocateInstance`, which skips all constructors and field initialisers. Ours gives safe defaults (`List.of()` etc.).

Q34. What happens if questions.json is missing or malformed?

The loader catches the exception, prints a [WARN], and returns `List.of()`. The app still starts; an interview would have 0 questions and a score of 0. This is tested in `loadUnknownResourceReturnsEmptyListInsteadOfThrowing`.

Q35. How do you add a new question? A new topic?

New question: add a JSON object to `questions.json` — no code change. New topic: add an enum constant (e.g. `STREAMS("Streams")`) and its questions; menus, balanced selection and reports pick it up automatically because they iterate `Topic.values()`.

Q36. Why is AnswerEvaluation immutable? Why List.copyOf?

A result shouldn't change after it is produced. Immutable objects are simpler, thread-safe, and safe to share with the UI and report. `List.copyOf` makes an unmodifiable defensive copy, so later changes to the caller's list don't leak in (and it rejects nulls).

Q37. Why does InMemoryQuestionBank return new ArrayList<>(questions)?

Defensive copy — callers (like setters that shuffle) can't modify the bank's internal list. Encapsulation of internal state.

Q38. Why is the best score initialised to -1?

A sentinel meaning "no interview yet", distinct from a genuine 0%. Alternatives: `Integer` with null, or `OptionalInt` — both more explicit; the sentinel is simple and is encapsulated behind a named constant `NO_SCORE_YET`.

Q39. Why use Map<Topic, int[]> in the report generator?

A compact mutable accumulator: index 0 = sum, index 1 = count, updated in place via `computeIfAbsent`. A small record/class (`TopicStats`) would be more readable; `EnumMap` would be faster and keep topic order.

11.5 Flow & control

Q40. Walk me through what happens when the user selects "1. Start mock interview".

The app asks for a strategy (Random/Balanced/Topic, 0 cancels) → creates `ConsoleInterviewUI` and `InterviewSession` → `run(ui)`: the setter selects 5 questions → `onSessionStart` → for each question: `onAskQuestion` reads the answer, `evaluator.evaluate` scores it, `onAnswerEvaluated` shows score/feedback/keywords/model answer → average → `candidate.recordInterviewResult` → `onSessionEnd` → the app prints the `PlacementReadinessReport` and returns to the menu (Fig. 5).

Q41. How do you handle invalid menu input like "abc"?

`readInt()` wraps `Integer.parseInt` in try/catch for `NumberFormatException` and returns `-1`, which falls into the switch's `default` branch ("Invalid option"). The program never crashes on bad input.

Q42. Why is a single Scanner shared across classes?

Multiple `Scanners` on `System.in` can buffer input and steal lines from each other; closing one closes `System.in` for all. So one `Scanner` is created in the app and passed (injected) to the UI classes, and it is closed once on exit.

Q43. Why `scanner.nextLine()` rather than `nextInt()`?

`nextInt()` leaves the newline in the buffer, causing the next `nextLine()` to return empty — a classic bug. Reading whole lines and parsing them avoids it and lets us handle bad input ourselves.

Q44. How many questions per session? Can it be changed?

Default `DEFAULT_QUESTION_COUNT = 5`; a second constructor accepts a custom count. If fewer questions are available (e.g. a topic with 3), setters return as many as exist (`Math.min(count, size)`).

11.6 Testing

Q45. How did you test the project?

34 JUnit 5 tests in 11 test classes, all passing (`mvn test`). Model: `CandidateTest` (3), `QuestionTest` (2), `TopicTest` (3). Service: `KeywordBasedEvaluatorTest` (5), `FallbackEvaluatorTest` (2), `InMemoryQuestionBankTest` (5), `Random-/Balanced-/TopicQuestionSetterTest` (3 each), `PlacementReadinessReportGeneratorTest` (3). Util: `JsonQuestionLoaderTest` (2). Plus a manual end-to-end demo run of the JAR (73%, ALMOST THERE).

Q46. How did you test the fallback logic without real AI or internet?

With test doubles: a nested `FailingAiEvaluator` whose `isAvailable()` returns true but `evaluate()` throws, and an `UnavailableAiEvaluator`. We assert the result comes from the keyword evaluator (score 100 for "finally") and that `isAiActive()` flips to false. This is possible only because `FallbackEvaluator` depends on an interface.

Q47. How do you test something random?

Inject a seeded `Random` (`new Random(5)`) through the constructor so output is repeatable, and assert *properties* (size, all from the requested topic) rather than exact order.

Q48. What was not tested?

The three AI clients end-to-end (no API key / local Ollama available during testing) and the console UI classes. Future work: mock the HTTP layer (e.g. `WireMock` or overriding the protected `send()` in `ClaudeLlmClient`) and test the UI with a `Scanner` over a `ByteArrayInputStream`.

11.7 Process & teamwork

Q49. Describe your development iterations.

Baseline: hard-coded questions, one KeywordBasedEvaluator, fixed random order. **Refinement:** extracted QuestionBank and AnswerEvaluator interfaces, moved questions to questions.json with Gson, added three question setters. **Final:** AI evaluators behind the same interface wrapped in FallbackEvaluator chosen by EvaluatorFactory, AI question generator, JUnit suite, shaded JAR.

Q50. Who did what in the team?

Rithika S built the service layer — the keyword and AI evaluators, FallbackEvaluator, EvaluatorFactory and the three question setters. Shwethaa JV built the data layer — the questions.json schema, JsonQuestionLoader, InMemoryQuestionBank — and wrote the 34 JUnit tests. We agreed on the shared interfaces first, so we could work in parallel (see §21.2).

Q51. What did you learn from this project?

Designing with interfaces first; how patterns (Strategy, Factory, Decorator) remove conditional logic; JSON handling and generics/type erasure; exception design and graceful degradation; calling REST APIs from Java; writing unit tests with test doubles; Maven builds; Git workflow.

Q52. What was the most challenging part?

Making AI optional yet seamless: handling network/parse failures without crashing, parsing semi-structured LLM output reliably, and keeping the engine unaware of all of it. The Decorator + custom exception design solved it.

Q53. If you had more time, what would you add?

Semantic/embedding-based scoring; persistent history (file or database via JDBC); a larger validated question bank; a GUI/web front-end on the existing service layer; difficulty levels and timed answers; end-to-end tests of AI evaluators; retry/circuit-breaker for AI calls.

12. Java language features used in the code

Examiners like to point at a line and ask "what is this?". Be able to explain each of these.

Feature (version)	Where	Explanation
Switch expression with <code>-></code> and <code>yield</code> (Java 14)	<code>App.run()</code> , <code>chooseQuestionSetter()</code>	Switch that returns a value; no fall-through, no <code>break</code> . <code>yield</code> returns a value from a block case.
Text blocks <code>"""</code> (Java 15)	<code>SYSTEM_PROMPT</code> in <code>AbstractLlmEvaluator</code> , <code>AiQuestionGenerator</code>	Multi-line string literals without <code>\n</code> and concatenation; incidental indentation removed.
Pattern matching for instanceof (Java 16)	<code>evaluator instanceof FallbackEvaluator fallback</code> in <code>EvaluatorFactory</code>	Tests and casts in one step, binding variable <code>fallback</code> .
<code>var</code> (Java 10)	<code>for (var q : questions) in browseQuestions()</code>	Local-variable type inference; still statically typed.
Lambdas & functional interfaces (Java 8)	<code>Topic deserializer</code> , <code>computeIfAbsent(topic, t -> new int[2])</code> , <code>forEach</code> , <code>test</code> , <code>allMatch(q -> ...)</code>	Anonymous function implementing a single-abstract-method interface.
Diamond with anonymous class (Java 9)	<code>new TypeToken<>() {}</code> in <code>JsonQuestionLoader</code>	Type inference for anonymous inner classes.
try-with-resources (Java 7)	<code>JsonQuestionLoader.load()</code>	Auto-closes <code>AutoCloseable</code> resources (<code>InputStream</code> , <code>Reader</code>) even on exceptions.
Multi-catch ordering	Gemini/Ollama clients	Specific catches (<code>IOException</code> , <code>InterruptedException</code> , <code>AiEvaluationException</code>) before generic <code>Exception</code> .
Immutable collection factories (Java 9/10)	<code>List.of()</code> , <code>List.copyOf()</code>	Unmodifiable lists; <code>copyOf</code> makes a defensive copy.
<code>Collections.unmodifiableList</code> , <code>shuffle</code>	<code>Question</code> , <code>setters</code> , <code>bank</code>	Read-only view; random permutation (Fisher–Yates) with optional <code>Random</code> .
String API (Java 11)	<code>isBlank()</code> , <code>repeat()</code> , <code>String.join</code> , <code>String.format/printf</code>	Blank check incl. whitespace; repeat for ASCII bars; join keyword lists.
<code>java.net.http.HttpClient</code> (Java 11)	Three LLM clients	Built-in HTTP/2 client with builders, timeouts, <code>BodyPublishers</code> / <code>BodyHandlers</code> .
Generics	<code>List<Question></code> , <code>Map<Topic,int[]></code> , <code>HttpResponse<String></code> , <code>JsonDeserializer<Topic></code>	Compile-time type safety; erased at runtime (hence <code>TypeToken</code>).
Enums with fields	<code>Topic</code>	<code>values()</code> , <code>valueOf()</code> , private constructor, comparison with <code>==</code> .
Annotations	<code>@Override</code> , <code>@SerializedName</code> , <code>@SuppressWarnings("unused")</code> , <code>@Test</code>	Metadata for compiler, Gson and JUnit.

Ternary & conditional add	<code>(cond ? matched : missed).add(k)</code>	Selects which list to add to in one expression.
StringBuilder	Report generator	Mutable, efficient string building in loops.
Environment variables	<code>System.getenv()</code> , <code>getOrDefault</code>	Externalised configuration.
Javadoc	Every class/method	<code>/** ... */</code> with <code>@param</code> , <code>@return</code> , <code>@author</code> ; generate with <code>mvn javadoc:javadoc</code> .

13. Core Java — OOP

Q1. What are the four pillars of OOP?

Encapsulation, Abstraction, Inheritance, Polymorphism (examples in chapter 7). This is literally question #2 in our question bank.

Q2. Abstract class vs interface?

Abstract class: can have state (instance fields), constructors, concrete and abstract methods, any access modifier; a class extends only one. Interface: no instance state; abstract, default, static and (Java 9) private methods; fields are public static final; a class can implement many → multiple inheritance of type. Use an interface for a capability/contract (AnswerEvaluator), an abstract class for shared code among related classes (AbstractLlmEvaluator).

Q3. Method overloading vs overriding?

Overloading: same name, different parameter list, same class; resolved at **compile time**. Overriding: subclass redefines an inherited method with the same signature; resolved at **runtime** (dynamic dispatch). Overriding can't reduce visibility or throw broader checked exceptions; return type may be covariant.

Q4. Can we override static or private methods?

No. Static methods are hidden, not overridden (resolved by reference type). Private methods aren't inherited. Final methods can't be overridden.

Q5. Why doesn't Java support multiple inheritance of classes?

To avoid the diamond problem (ambiguity when two parents have the same method/state). Java allows multiple interfaces; if two default methods conflict, the class must override and choose (A.super.m()).

Q6. What is a constructor? Can it be private, static, final or inherited?

Special method that initialises a new object; same name as class, no return type. It can be private (Singleton, utility classes like our EvaluatorFactory). It cannot be static, final, abstract, or inherited. this() calls another constructor of the same class; super() calls the parent's — must be the first statement.

Q7. What does the final keyword mean?

final variable — assigned once (constant reference, object may still be mutable); final method — can't be overridden; final class — can't be extended (String, our EvaluatorFactory, JsonQuestionLoader).

Q8. What does the static keyword mean?

Belongs to the class, not instances: static fields (shared), static methods (no this, can't access instance members directly), static blocks (run once on class load), static nested classes (no outer instance — used for our test stubs).

Q9. Access modifiers?

private (class only) < default/package-private (same package) < protected (package + subclasses) < public (everywhere).

Q10. What is composition vs inheritance? Which is better?

Inheritance = "is-a" (ClaudeAnswerEvaluator is-a AbstractLlmEvaluator). Composition = "has-a" (BalancedQuestionSetter has-a RandomQuestionSetter). Prefer composition: looser coupling, change behaviour at runtime, avoids fragile base-class problems.

Q11. equals() and hashCode() contract?

If `a.equals(b)` then `a.hashCode() == b.hashCode()`. Unequal objects may share a hash. Override both together, or `HashMap/HashSet` break. `==` compares references; `equals` compares logical equality. (Our enum `Topic` is compared with `==`, which is correct for enums.)

Q12. What is an immutable class and how do you create one?

State cannot change after construction. Make the class `final`, fields `private final`, no setters, initialise in constructor, make defensive copies of mutable inputs/outputs. Examples: `String`, wrapper classes, our `AnswerEvaluation`.

Q13. What is the Object class? Its key methods?

Root of all classes. `equals`, `hashCode`, `toString`, `getClass`, `clone`, `finalize` (deprecated), `wait`, `notify`, `notifyAll`. We override `toString()` in `Candidate`.

Q14. What is an enum? Can it have constructors and methods?

A type-safe set of constants; each constant is a singleton instance. Yes — private constructors, fields, methods, and it can implement interfaces (see `Topic`). Gets `values()`, `valueOf()`, `ordinal()`, `name()`. Works in switch statements.

Q15. What are inner, static nested, local and anonymous classes?

Inner (non-static) needs an outer instance; static nested doesn't; local is declared in a method; anonymous is an unnamed one-off subclass/implementation (our `new TokenType<...>(){}()`).

Q16. What is a marker interface? Functional interface?

Marker: no methods, tags a capability (`Serializable`, `Cloneable`). Functional: exactly one abstract method, can be a lambda target (`Runnable`, `Comparator`, `Function`; Gson's `JsonDeserializer` is effectively one).

Q17. Shallow vs deep copy?

Shallow copies field references (shared inner objects); deep copies inner objects recursively. `new ArrayList<>(list)` is a shallow copy — the `Question` objects are shared, which is fine because `Question` is effectively immutable.

Q18. Pass by value or reference?

Java is always pass-by-value. For objects, the *reference* is copied, so the method can mutate the object but can't make the caller's variable point elsewhere.

14. Core Java — Collections

Q1. Explain the Collection hierarchy.

Iterable → Collection → List (ArrayList, LinkedList, Vector), Set (HashSet, LinkedHashSet, TreeSet), Queue / Deque (PriorityQueue, ArrayDeque, LinkedList). Map is separate (HashMap, LinkedHashMap, TreeMap, Hashtable, ConcurrentHashMap, EnumMap).

Q2. ArrayList vs LinkedList?

ArrayList: dynamic array, O(1) random access, amortised O(1) append, O(n) insert/remove in the middle, cache-friendly. LinkedList: doubly linked, O(n) access, O(1) insert/remove at known node, more memory. We use ArrayList everywhere because we mostly append and index (BalancedQuestionSetter uses `pool.get(cursor[i])`).

Q3. How does HashMap work internally?

Array of buckets (default 16, load factor 0.75). `hash(key)` → index; collisions stored in a linked list, converted to a red-black tree when a bucket exceeds 8 entries (Java 8). Resizes (doubles) when `size > capacity × load factor`. One null key allowed. Average O(1) get/put. Not synchronized; iteration order not guaranteed.

Q4. HashMap vs Hashtable vs ConcurrentHashMap?

Hashtable: legacy, every method synchronized, no null keys/values. HashMap: not thread-safe, allows one null key. ConcurrentHashMap: thread-safe with fine-grained locking/CAS, no nulls, high concurrency.

Q5. HashMap vs LinkedHashMap vs TreeMap vs EnumMap?

HashMap — no order; LinkedHashMap — insertion (or access) order, used for LRU caches; TreeMap — sorted by key, O(log n), red-black tree; EnumMap — enum keys in an array, ordered by ordinal, very fast. *In our report, EnumMap would print topics in enum order.*

Q6. List vs Set vs Map?

List: ordered, allows duplicates, index access. Set: no duplicates. Map: key→value pairs, unique keys.

Q7. HashSet internals?

Backed by a HashMap; elements are keys with a dummy constant value. Relies on `hashCode/equals`.

Q8. Comparable vs Comparator?

Comparable: natural ordering inside the class, `compareTo(o)`, one ordering. Comparator: external, `compare(a,b)`, many orderings, lambda-friendly (`Comparator.comparing(Question::getId)`).

Q9. Fail-fast vs fail-safe iterators?

Fail-fast (ArrayList, HashMap) throw `ConcurrentModificationException` if the collection is structurally modified during iteration (tracked via `modCount`). Fail-safe (CopyOnWriteArrayList, ConcurrentHashMap) iterate over a snapshot/weakly consistent view. Use `Iterator.remove()` or `removeIf` to remove safely.

Q10. Unmodifiable vs immutable collections?

`Collections.unmodifiableList(list)` is a read-only view — changes to the underlying list are visible. `List.of` / `List.copyOf` are truly immutable and reject nulls. Both throw `UnsupportedOperationException` on mutation (tested in `expectedKeywordsAreUnmodifiable`).

Q11. What does Collections.shuffle do?

Random permutation in O(n) using the Fisher–Yates algorithm; overload takes a `Random` for reproducibility (we inject a seeded `Random` in tests).

Q12. What does `subList` return?

A *view* backed by the original list, not a copy. Structural changes to the original invalidate it. That's why `InMemoryQuestionBank.getRandomQuestions` wraps it: `new ArrayList<>(copy.subList(...))`.

Q13. What is `computeIfAbsent`?

Java 8 `Map` method: if the key is missing, computes a value with the given function, stores and returns it; otherwise returns the existing value. Replaces the `get/null-check/put` idiom (used in our report generator).

Q14. Array vs ArrayList?

`Array`: fixed size, can hold primitives, `length` field. `ArrayList`: resizable, objects only (autoboxing), rich API, `size()`. We use an `int[]` for cursors and per-topic sums because they're fixed-size primitives.

Q15. What is a `PriorityQueue`? `Deque`?

`PriorityQueue`: binary heap, `poll` returns min (or by comparator), $O(\log n)$. `Deque`: double-ended queue (`ArrayDeque`) — use as stack or queue; preferred over `Stack/LinkedList`.

15. Core Java — Strings

Q1. Why is String immutable?

Security (class names, file paths, URLs can't be changed after checks), safe sharing in the String Pool, thread-safety, and cached hashCode (great as HashMap keys).

Q2. What is the String Constant Pool?

A special area in the heap storing unique string literals. "abc" literals are reused; new String("abc") creates a new heap object. intern() returns the pooled instance.

Q3. String vs StringBuilder vs StringBuffer?

String — immutable. StringBuilder — mutable, not synchronized, fastest (used in our report generator). StringBuffer — mutable, synchronized (thread-safe, slower). Use StringBuilder for concatenation in loops.

Q4. == vs equals() for strings?

== compares references; equals() compares characters. "a" == "a" is true (pool) but new String("a") == "a" is false. Use equals / equalsIgnoreCase.

Q5. How many objects does new String("hello") create?

Up to two: one literal in the pool (if not already present) and one new object on the heap.

Q6. isEmpty() vs isBlank()?

isEmpty() — length is 0. isBlank() (Java 11) — empty or only whitespace. We use isBlank() to detect skipped answers.

Q7. How do you do case-insensitive matching like your evaluator?

Normalise both sides with toLowerCase() then contains(); Or equalsIgnoreCase for whole strings; or a regex with Pattern.CASE_INSENSITIVE and \b word boundaries (better — avoids "final" matching "finally").

Q8. Useful String methods?

substring, indexOf, lastIndexOf (used in extractJson), trim/strip, split, join, replace, format, repeat, chars, toCharArray, valueOf, matches.

Q9. Why is char[] preferred over String for passwords?

Strings are immutable and may linger in the pool/heap until GC; a char[] can be wiped after use.

16. Core Java — Exceptions

Q1. Exception hierarchy?

`Throwable` → `Error` (`OutOfMemoryError`, `StackOverflowError` — don't catch) and `Exception` → checked exceptions (`IOException`, `InterruptedException`) and `RuntimeException` (unchecked: `NullPointerException`, `IllegalArgumentException`, `NumberFormatException`, our `AiEvaluationException`).

Q2. Checked vs unchecked exceptions?

Checked: verified at compile time; must be caught or declared with `throws` (`IOException` from `HttpClient.send`).

Unchecked: subclasses of `RuntimeException`, not enforced — usually programming errors or conditions callers can't meaningfully handle at every level.

Q3. throw vs throws?

`throw` raises an exception object (`throw new AiEvaluationException(...)`). `throws` in a method signature declares possible checked exceptions.

Q4. final vs finally vs finalize?

`final` — keyword for constants/no override/no subclass. `finally` — block that always runs after try/catch (cleanup). `finalize()` — deprecated GC hook; don't use.

Q5. Does finally always execute?

Almost always — even after `return`. Exceptions: `System.exit()`, JVM crash, thread killed, infinite loop in try.

Q6. What is try-with-resources?

Java 7 construct that auto-closes any `AutoCloseable` declared in the parentheses, in reverse order, even if an exception occurs; exceptions from `close()` become *suppressed*. Used in `JsonQuestionLoader` for the `InputStream` and `Reader`.

Q7. How do you create a custom exception?

Extend `Exception` (checked) or `RuntimeException` (unchecked), provide constructors with message and cause — exactly our `AiEvaluationException(String)` and `(String, Throwable)`.

Q8. What is exception chaining?

Wrapping a low-level exception as the *cause* of a higher-level one (`new AiEvaluationException("Network error", e)`) so the original stack trace isn't lost; accessible via `getCause()`.

Q9. Order of catch blocks?

Most specific first; a superclass catch before a subclass catch is a compile error (unreachable). Our Gemini client catches `IOException`, `InterruptedException`, `AiEvaluationException` (rethrow) before generic `Exception`. Multi-catch: `catch (IOException | ParseException e)`.

Q10. Is catching Exception/Runtime Exception bad practice?

Generally yes, as it can hide bugs. It's acceptable at boundaries where any failure must be handled uniformly — like `FallbackEvaluator`, which must keep the interview running whatever the AI does — and we log the message.

Q11. What is NumberFormatException and where do you handle it?

Unchecked exception from `Integer.parseInt` on non-numeric text; caught in `readInt()`, which returns `-1`.

17. Core Java — JVM & memory

Q1. JDK vs JRE vs JVM?

JVM executes bytecode (platform-specific). JRE = JVM + core libraries (to run). JDK = JRE + tools (javac, jar, javadoc, jshell) to develop. Since Java 11 there's no separate JRE download.

Q2. How is Java platform-independent?

javac compiles source to platform-neutral bytecode (.class); each OS has its own JVM that interprets/JIT-compiles it — “write once, run anywhere”. Our JAR runs unchanged on Windows/Linux/Mac.

Q3. JVM memory areas?

Heap (objects, shared; young gen Eden + survivors, old gen), **Stack** (per thread: frames with local variables and references), **Metaspace** (class metadata, replaced PermGen in Java 8), **PC register**, **native method stack**. String pool lives in the heap.

Q4. Stack vs heap?

Stack: per-thread, LIFO, primitives and references, auto-freed when the method returns, StackOverflowError on deep recursion. Heap: shared, all objects, managed by GC, OutOfMemoryError when full.

Q5. What is garbage collection? How does it work?

Automatic reclamation of unreachable objects. Generational: most objects die young → frequent minor GCs in young gen; survivors are promoted to old gen → major GC. Mark (reachable from GC roots) → sweep/compact. Collectors: Serial, Parallel, G1 (default since 9), ZGC, Shenandoah. `System.gc()` is only a hint.

Q6. Class loading process & class loaders?

Loading → Linking (verify, prepare, resolve) → Initialisation (static blocks). Loaders: Bootstrap (core java.*), Platform/Extension, Application/System (classpath — loads our classes and finds questions.json). Parent-delegation model: ask the parent first.

Q7. What is JIT?

Just-In-Time compiler in the JVM (HotSpot C1/C2) compiles frequently executed bytecode (“hot spots”) to native machine code at runtime for speed.

Q8. What is the classpath?

The list of directories/JARs where the JVM looks for classes and resources. In the shaded JAR, Gson classes and questions.json are on the classpath inside the JAR.

Q9. Memory leaks in Java — possible?

Yes — objects still referenced but no longer needed (static collections that grow, unclosed resources, listeners never removed, ThreadLocals). Our question bank grows with AI-generated questions, but it's bounded by user actions.

Q10. Strong, weak, soft, phantom references?

Strong: normal, never collected while reachable. Soft: collected when memory is low (caches). Weak: collected at next GC (WeakHashMap). Phantom: post-mortem cleanup tracking.

18. Core Java — Multithreading

Q1. Is your application multithreaded?

No — it's single-threaded by design: an interview is a sequential conversation. We use the *synchronous* `HttpClient.send()`. We do handle `InterruptedException` correctly. If we added a GUI, the AI call should move off the UI thread (e.g. `sendAsync()` returning a `CompletableFuture`, or a `SwingWorker`).

Q2. Process vs thread?

Process: independent program with its own memory. Thread: lightweight unit of execution inside a process sharing heap memory; each has its own stack.

Q3. Ways to create a thread?

Extend `Thread` and override `run()`; implement `Runnable` (preferred — keeps inheritance free); implement `Callable` (returns a value, can throw) with an `ExecutorService`; Java 21 virtual threads.

Q4. `start()` vs `run()`?

`start()` creates a new thread that calls `run()`. Calling `run()` directly just executes it on the current thread.

Q5. Thread life-cycle?

NEW → RUNNABLE → (BLOCKED / WAITING / TIMED_WAITING) → TERMINATED.

Q6. What is synchronization?

Ensures only one thread executes a critical section on a monitor at a time (`synchronized` methods/blocks), preventing race conditions and providing visibility (`happens-before`).

Q7. `volatile`?

Guarantees visibility of writes across threads and prevents reordering, but not atomicity (`count++` is still unsafe → use `AtomicInteger`).

Q8. `wait/notify` vs `sleep`?

`wait()` releases the monitor and waits for `notify()`; must be called in synchronized code; Object method. `sleep()` pauses without releasing locks; Thread static method.

Q9. What is a deadlock? How to avoid it?

Two+ threads each hold a lock the other needs, waiting forever. Avoid with consistent lock ordering, timeouts (`tryLock`), fewer nested locks.

Q10. `ExecutorService` and thread pools?

Framework to reuse a pool of threads: `Executors.newFixedThreadPool(n)`, `submit()` returns a `Future`; `shutdown()` when done. Avoids the cost of creating threads per task.

Q11. Is your code thread-safe?

Not required. `AnswerEvaluation` and `Question` are effectively immutable and thus thread-safe; `Candidate`, `InMemoryQuestionBank` (`ArrayList`) and `FallbackEvaluator.aiActive` are mutable and not synchronized. To parallelise we'd use `CopyOnWriteArrayList`, `AtomicInteger`, `volatile` / `AtomicBoolean` for `aiActive`.

Q12. What is interruption?

A cooperative cancellation signal: `thread.interrupt()` sets a flag; blocking methods throw `InterruptedException` and clear it — so we restore it with `Thread.currentThread().interrupt()`.

19. Java 8+ features

Q1. Main Java 8 features?

Lambdas, functional interfaces, Stream API, method references, default and static interface methods, `Optional`, new Date/Time API (`java.time`), `CompletableFuture`, Map enhancements (`computeIfAbsent`, `getOrDefault`, `forEach`).

Q2. What is a lambda expression?

A concise anonymous function implementing a functional interface: `(params) -> expression`. Example from our code: `t -> new int[2]` and the Gson deserializer lambda.

Q3. What is the Stream API? Intermediate vs terminal operations?

Declarative processing of collections. Intermediate (lazy, return a stream): `filter`, `map`, `sorted`, `distinct`, `limit`. Terminal (trigger execution): `collect`, `forEach`, `count`, `reduce`, `anyMatch/allMatch`. Our `getQuestionsByTopic` could be written as `questions.stream().filter(q -> q.getTopic() == topic).toList()`.

Q4. map vs flatMap?

`map` transforms each element one-to-one; `flatMap` maps each element to a stream and flattens the results (e.g. all keywords of all questions).

Q5. What is Optional?

A container that may or may not hold a value, to avoid null checks and NPEs: `Optional.ofNullable(x).orElse(default)`. We could have used it in `chooseTopic()` instead of returning null.

Q6. Method references?

Shorthand for lambdas: `Topic::getDisplayName`, `System.out::println`, `ArrayList::new`.

Q7. Default methods in interfaces — why?

To add new methods to existing interfaces without breaking implementations (e.g. `List.sort`). We could add default `String describe()` to `AnswerEvaluator`.

Q8. What are records (Java 16)? Would you use them?

Concise immutable data carriers: `record AnswerEvaluation(Question question, int score, ...)` auto-generates constructor, accessors, equals, hashCode, toString. `AnswerEvaluation` and a `TopicStats` helper would be good candidates. (Gson 2.10+ supports records.)

Q9. Sealed classes (Java 17)?

sealed interface `AnswerEvaluator` permits `KeywordBasedEvaluator`, `FallbackEvaluator`, ... restricts which classes may implement a type — useful for exhaustive switch pattern matching.

Q10. Switch expressions vs old switch?

Arrow labels, no fall-through, can return values, compiler checks exhaustiveness for enums/sealed types. We use it in the main menu and strategy chooser.

20. Testing (JUnit 5), Maven & Git

20.1 Our test suite

Test class	#	Test methods
CandidateTest	3	newCandidateHasNoScoreYet · recordInterviewResultTracksBestScoreAndCount · toStringIncludesLatestBestScore
QuestionTest	2	gettersReturnConstructorValues · expectedKeywordsAreUnmodifiable
TopicTest	3	hasSixCoreTopics · displayNameAreHumanReadable · valueOfResolvesEnumConstantFromName
KeywordBasedEvaluatorTest	5	fullKeywordMatchScoresHundred · partialKeywordMatchScoresProportionally · blankAnswerScoresZeroWithNoAnswerFeedback · nullAnswersTreatedAsNoMatches · keywordMatchingIsCaseInsensitive
FallbackEvaluatorTest	2	fallsBackToStaticEvaluatorWhenAiThrows · usesStaticEvaluatorImmediatelyWhenAiUnavailable
InMemoryQuestionBankTest	5	loadsQuestionsFromClasspathJsonOnConstruction · getQuestionsByTopicFiltersCorrectly · getRandomQuestionsNeverExceedsRequestedCount · getRandomQuestionsCapsAtBankSizeWhenCountIsTooLarge · addQuestionsAppendsToTheBank
RandomQuestionSetterTest	3	selectQuestionsReturnsRequestedCount · selectQuestionsByTopicOnlyReturnsThatTopic · strategyNameIsRandom
BalancedQuestionSetterTest	3	distributesOneQuestionPerTopicBeforeRepeating · topicScopedSelectionDelegatesToSingleTopic · strategyNameMentionsAllTopics
TopicQuestionSetterTest	3	onlyReturnsQuestionsFromBoundTopic · defaultConstructorBindsToOopTopic · strategyNameIncludesTopicDisplayName
JsonQuestionLoaderTest	2	loadDefaultReturnsAllQuestionsFromResourceFile · loadUnknownResourceReturnsEmptyListInsteadOfThrowing
PlacementReadinessReportGeneratorTest	3	highAverage...Ready · lowAverage...NeedsPreparation · midRange...AlmostThere
Total: 11 classes	34	34 passed, 0 failures

20.2 Testing questions

Q1. What is unit testing? What is JUnit?

Testing the smallest units (methods/classes) in isolation. JUnit 5 (Jupiter) is the standard Java framework: `@Test`, assertions (`assertEquals`, `assertTrue`, `assertThrows`), lifecycle (`@BeforeEach`, `@AfterEach`, `@BeforeAll`), `@DisplayName`, `@ParameterizedTest`.

Q2. JUnit 4 vs JUnit 5?

JUnit 5 = Platform + Jupiter + Vintage; package `org.junit.jupiter.api`; `@BeforeEach` replaces `@Before`; test classes/methods needn't be public; `assertThrows` replaces `@Test(expected=...)`; supports lambdas, nested and parameterised tests.

Q3. Unit vs integration vs system testing?

Unit: one class in isolation (our 34 tests). Integration: components together (e.g. `InMemoryQuestionBank` with the real JSON file — partly covered). System/end-to-end: whole app (our manual JAR demo run). Also black-box vs white-box.

Q4. What are mocks / stubs / test doubles?

Stand-ins for real dependencies. Stub returns canned data; mock also verifies interactions (Mockito). We hand-wrote stubs (`FailingAiEvaluator`) — possible because we depend on interfaces.

Q5. What is code coverage?

Percentage of code executed by tests (line/branch). Measured with JaCoCo. High coverage $\neq$ correct tests, but low coverage shows untested risk.

Q6. What is TDD?

Test-Driven Development: Red (write a failing test) → Green (minimal code to pass) → Refactor.

20.3 Maven & Git questions

Q7. What is Maven and pom.xml?

A build & dependency-management tool. `pom.xml` (Project Object Model) declares coordinates (groupId `com.interviewsim`, artifactId `interview-simulator`, version `1.0.0`), dependencies with scopes, properties (Java 17) and plugins. Dependencies are downloaded from Maven Central into `~/.m2`.

Q8. What does the dependency scope `test` mean?

JUnit is available only when compiling/running tests, not packaged into the final JAR. Default scope `compile` (Gson) is available everywhere and is bundled by shade.

Q9. Fat JAR vs normal JAR?

A normal JAR contains only our classes; dependencies must be on the classpath separately. A fat/uber JAR (maven-shade-plugin) bundles dependencies, and we set `Main-Class` via `ManifestResourceTransformer` so `java -jar` works.

Q10. What Git workflow did you use?

Commits per milestone and feature branches (e.g. `poster-and-report-update`) merged into `main`. Commands: `git add`, `commit`, `branch`, `checkout/switch`, `merge`, `push`, `pull`, `log`, `status`. `.gitignore` keeps `target/` out of the repo.

21. The PBL report: walkthrough & questions

Examiners usually have the report open while asking questions. This chapter follows the report chapter by chapter, with the facts to quote and the questions each chapter tends to trigger.

21.1 Front matter

Item	What the report says
Title	AN INTERACTIVE JAVA TECHNICAL REVIEW SIMULATOR (the application itself is called “Java AI Interview Simulator”)
Team	RITHIKA S (2104251040779) · SHWETHAA JV (2104251050318)
Supervisor	Dr. R. Kavitha, M.E., Ph.D., Associate Professor, Dept. of CSE (also thanked as Project Co-ordinator)
Head of Department	Dr. S. Pavithra, M.E., Ph.D.
Others acknowledged	Chairman Shri P. Sriram · Principal Dr. A. Ramesh · Dean Dr. V. Srinivasa Rao · Class advisors Sankari C. and Sri Sai Sanjai
Keywords	Java, Object-Oriented Programming, Interview Simulation, JUnit Testing, Strategy Pattern
Repository	github.com/Rithikas07/interview-simulator (Appendix A.1)

Q1. Why is the report titled “Interactive Java Technical Review Simulator” while the app says “Java AI Interview Simulator”?

“Technical review” is the formal title for the PBL; the app name describes what it does — an AI-assisted mock interview. Both refer to the same Maven project (`com.interviewsim`).

21.2 Who did what (Chapter 7.1) — be ready to defend your own part

RITHIKA S — service layer

- KeywordBasedEvaluator
- AiAnswerEvaluator implementations (Claude, Gemini, Ollama)
- FallbackEvaluator (Decorator) and EvaluatorFactory
- Random-, Balanced-, TopicQuestionSetter (Strategy)

Expect: Strategy vs Decorator, the fallback flow (Fig. 7), the round-robin algorithm (§10.2), the AI prompt and JSON parsing (§10.4), HttpClient, why `AiEvaluationException` is unchecked, the Bridge between evaluators and `LlmClient`.

SHWETHAA JV — data layer & tests

- questions.json schema (id, topic, question, keywords, modelAnswer)
- JsonQuestionLoader (Gson, custom Topic enum deserializer)
- InMemoryQuestionBank
- The 34 JUnit 5 tests across model, service and util

Expect: `TypeToken` and type erasure, `@SerializedName`, `getResourceAsStream`, `try-with-resources`, what happens with a missing file, seeded `Random` in tests, test doubles in `FallbackEvaluatorTest`, JUnit 5 annotations and assertions (§20).

Important: examiners often ask each member about the *other* member’s module to check that both understand the whole system. Both of you should be able to explain the architecture (Fig. 1), the interview sequence (Fig. 5) and all four main patterns.

Q2. How did you divide the work, and how did you integrate it? (Chapter 7.2)

By architectural layer — Rithika on evaluators/strategies, Shwethaa on the data layer and tests. We first agreed on the shared interfaces (`AnswerEvaluator`, `QuestionBank`, `QuestionSetter`), so each of us could build and test independently and integrate late without major conflicts. Self and peer assessment: 50% / 50% (Table A.1).

Q3. What feedback did your mentor give, and how did you act on it?

Early feedback was against hard-coding — which led to moving questions into `questions.json` and putting the bank and evaluator behind interfaces. Later remarks asked for unit tests before integration, which led to the JUnit suite (Table 3.1).

21.3 Chapter 1 — Introduction

- **Background:** placement drives include a core-Java round; students practise with static notes; no interactive tool gives instant, topic-wise feedback. Building the simulator in Java reinforces the same concepts it tests.
- **Requirements of a good solution:** runs offline, gives topic-wise diagnostic feedback, small enough for two students in one PBL cycle.
- **Driving question:** see §1.3 — memorise it word for word.
- **Scope & limitations (1.4):** single user, console, no persistent storage, 23 bundled questions, AI needs an API key or local Ollama.

Q4. What is a “driving question” in PBL?

The open-ended question that frames the whole project and that the final product must answer. Ours asks whether a console Java app can run a realistic topic-focused mock interview, score it automatically and report readiness without a network or paid service — and Chapter 8 concludes that it can.

21.4 Chapter 2 — Concept exploration

Approach	Result	Decision
Core Java + Collections + Gson, console, Maven + shade (refs [1],[2],[3],[6])	Sub-second start-up, no external services, self-contained JAR	Chosen
Strategy / Factory / Decorator patterns (ref [5] GoF)	3 interchangeable question setters and a fallback evaluator chain without conditionals in the entry point	Adopted
Spring Boot + JPA/Hibernate + REST	Multi-user web UI and persistent history, but much more setup	Considered, rejected for scope

Q5. Why didn't you use Spring Boot?

The PBL's learning objectives are core Java. Spring Boot would add an embedded server, ORM configuration and a front-end — features we didn't need — and delay the first runnable demo. Our service layer is UI-agnostic, so a Spring REST layer could be added later.

21.5 Chapter 3 — Planning

Weeks	Milestone	Work done
1–2	Problem statement, objectives, time plan	Driving question; chose core Java + Gson
3–4	Data handling, methodology design	Question/Topic/Candidate model and the <code>questions.json</code> schema
5–7	Implementation & testing	<code>QuestionBank</code> , <code>QuestionSetter</code> strategies, evaluator chain, report generator, JUnit 5 suite
8–10	Presentation with demo	Executable JAR and scripted mock-interview demo

11–12	Final testing & documentation	34/34 tests passing; report and poster
-------	-------------------------------	--

Requirements (Table 3.2): Intel i5 or equivalent, 8 GB RAM, Windows 11 · Java 17 (LTS) · Gson 2.11.0, JUnit Jupiter 5.10.2 · Apache Maven 3.9.6 (compiler, surefire, shade plugins) · no database · VS Code / IntelliJ IDEA · Git.

Q6. How did you make sure the project was feasible?

We kept the scope small (console, no database, no UI framework) and made every milestone a runnable increment, so by the implementation weeks we were extending a working vertical slice rather than integrating everything at the end.

21.6 Chapter 4 — Iterative design

Iteration	What it had	Problem that led to the next step
Baseline	Hard-coded questions, one KeywordBasedEvaluator, fixed random order	Rigid — questions and scoring couldn’t be swapped
Refinement	QuestionBank & AnswerEvaluator interfaces, questions.json via Gson, three QuestionSetter strategies	Keyword scoring misses correct but differently-worded answers
Final	AI evaluators behind the same interface, wrapped in FallbackEvaluator, chosen by EvaluatorFactory; package separation model / service / core / ui / util	—

Q7. Explain Figure 4.1 (system architecture).

Registration → strategy selection → question bank (JSON via Gson) → answer evaluation (AI + fallback / keyword) → readiness report. CandidateRegistrationUI builds the Candidate; EvaluatorFactory picks the evaluator (Claude → Gemini → Ollama → keywords); InterviewSession pulls questions from InMemoryQuestionBank using the chosen QuestionSetter, asks through ConsoleInterviewUI, evaluates, and PlacementReadinessReportGenerator prints the report. (Use Fig. 1 and Fig. 5 of this guide.)

21.7 Chapter 5 — Implementation

Module (5.1)	Classes
Input	CandidateRegistrationUI, ConsoleInterviewUI
Processing	InterviewSession, QuestionSetter implementations
Data management	InMemoryQuestionBank, JsonQuestionLoader, questions.json
Validation / evaluation	KeywordBasedEvaluator, AI evaluators, FallbackEvaluator, EvaluatorFactory
Output	PlacementReadinessReportGenerator

The three snippets in 5.2 are FallbackEvaluator.evaluate() (§6.3), the keyword-scoring loop (§10.1) and the round-robin loop (§10.2). Be ready to explain each line. Figures 5.1 and 5.2 are screenshots of the input (registration, menu, a typed answer) and output (evaluation and final report).

21.8 Chapter 6 — Results

Metrics (6.1): (a) JUnit pass rate — correctness of model, evaluator and selection logic; (b) the app’s own keyword-match percentage for a sample session — its placement-readiness metric. Code coverage was not measured.

Sample session (Balanced strategy, keyword evaluator)	OOP	Collections	Strings	Exceptions	JVM	Overall
Score	83%	100%	83%	0%	100%	73% — ALMOST THERE

Q8. Check the arithmetic: how is the overall 73%?

$(83 + 100 + 83 + 0 + 100) \div 5 = 366 \div 5 = 73.2 \rightarrow \text{Math.round} \rightarrow 73\%$. 73 is ≥ 50 but < 75 , so the verdict is ALMOST THERE.

Q9. Why did Exceptions score 0%?

The question drawn was about try-with-resources, but the prepared answer was about checked vs unchecked exceptions, so none of the expected keywords matched. That is the main limitation of keyword matching: it can't tell a related-but-different answer, or a correct answer worded differently, from a wrong one. The AI evaluators were added to address this but weren't exercised in the run (no API key or local Ollama).

Q10. Why only five topics in the sample — where is Multithreading?

The session had 5 questions and BalancedQuestionSetter takes one question per topic in enum order: OOP, Collections, Strings, Exceptions, JVM. Multithreading would come sixth. This shows the round-robin working as designed (6.3).

Q11. What limitations did you report? (6.4)

Keyword scoring can't judge correct answers that lack the exact keywords; no concurrency or multi-user testing; no code-coverage tool; AI evaluators not tested end-to-end.

21.9 Chapter 7 — Course outcomes

Outcome	Evidence
CO1 — Technical skill	Pluggable evaluator chain and Strategy-pattern question setters (Chapters 4, 5)
CO2 — Testing discipline	34 JUnit 5 tests passing across model, service and util (Table 6.1)
CO3 — Teamwork & process	Weekly milestone log (Table 3.1) showing incremental delivery

21.10 Chapter 8 — Conclusion, future scope & references

Conclusion: the app answers the driving question — offline-capable, 23-question JSON bank, three question strategies, resilient evaluator chain, topic-wise report, 34/34 tests.

Future scope: semantic/embedding scoring · persistent storage · JaCoCo coverage · bigger, validated question bank · GUI or web front-end on the existing service layer · end-to-end testing of the AI evaluators.

References: [1] Oracle Java Tutorials · [2] Java SE 17 API Specification · [3] Gson User Guide · [4] JUnit 5 User Guide · [5] Gamma et al., *Design Patterns* (1994) · [6] Java Collections Framework Overview · [7] Maven Shade Plugin documentation.

Q12. What is JaCoCo and why is it in your future scope?

Java Code Coverage — a tool that plugs into Maven and reports which lines and branches the tests execute. We only reported pass/fail, so adding JaCoCo would show how much of the code the 34 tests actually exercise.

22. Limitations, tricky questions & improvements

Examiners often look for weaknesses. Acknowledging them calmly and proposing a fix is far better than being caught out. These are real observations from the current code.

Observation	Why it matters	How you would fix it
Keyword scoring uses <code>contains</code>	"final" matches "finally"; synonyms and paraphrases get no credit; keyword stuffing scores 100.	Word-boundary regex, synonyms per keyword, stemming; AI or embedding similarity; penalise very short answers.
No persistence	Candidate history and AI-generated questions are lost on exit.	Save to a JSON file with Gson, or a database via JDBC/SQLite; profiles keyed by roll number.
<code>describeSelectedEvaluator()</code> prints <code>FallbackEvaluator</code>	It uses <code>fallback.getClass()</code> — the wrapper's class — so the banner can't show "Claude/Gemini/Ollama".	Add a getter for the primary evaluator (or provider name) and print that.
<code>selectBestEvaluator()</code> is called twice at startup	The second call (via <code>describeSelectedEvaluator</code>) re-creates evaluators and re-probes Ollama over HTTP (up to 3 s).	Describe the evaluator already selected: pass it as a parameter.
AI fallback is permanent	One transient error disables AI for the whole run.	Retry with back-off; circuit breaker that retries after N questions.
AI question IDs always start at 1000	A new <code>AiQuestionGenerator(client, 1000)</code> is created per request, so repeated generation produces duplicate ids.	Keep one generator, or compute next id from the bank's max id.
<code>findClientFor()</code> uses an <code>instanceof</code> chain	Violates Open/Closed — adding a provider means editing the app.	Expose <code>getClient()</code> on <code>AiAnswerEvaluator</code> or let the factory provide clients.
Report topic order uses <code>HashMap</code>	Topic lines may appear in an unpredictable order.	Use <code>EnumMap<Topic, ...></code> (ordered, faster).
"Overall score ... : 365/5 (73%)" format	The first number is the sum of percentages over the number of questions — can confuse readers.	Print "73% average over 5 questions".
Report generator takes Candidate twice	Constructor stores a candidate that <code>generateReport</code> doesn't use (it uses the parameter).	Remove the unused field/constructor parameter.
Gemini key in URL query	URLs can end up in logs/proxies.	Send it in the <code>x-goog-api-key</code> header.
Small bank (23 Qs), uneven topics	Repetition after a few sessions; topic averages based on 1 question.	Grow and validate the bank; AI generation with review; difficulty tags.
AI evaluators not tested end-to-end	No API key / local Ollama during testing.	Mock HTTP (override protected <code>send()</code> , <code>WireMock</code>); contract tests on sample JSON replies.

22.1 Tricky questions

Q1. Is FallbackEvaluator really a Decorator? It wraps two objects.

It follows the Decorator structure — same interface, wraps a component, adds behaviour (fault tolerance) transparently. Because it chooses between two delegates, it also resembles a Proxy and a 2-step Chain of Responsibility. The intent — adding resilience without changing callers — is the decorator intent.

Q2. Is EvaluatorFactory the GoF Factory Method?

No — it's a Simple (static) Factory. GoF Factory Method defines an abstract creation method that subclasses override. Abstract Factory creates families of related objects. Ours is a static method that chooses and returns one product via its interface.

Q3. Can AI scoring be trusted? It may give different scores for the same answer.

LLMs are non-deterministic. We reduce variance with a strict rubric, low temperature (0.4 for Gemini/Ollama), JSON-only replies and clamping, and we still show matched/missed keywords. For high-stakes use we'd average multiple runs or combine AI and keyword scores.

Q4. What if the candidate types an answer that tries to manipulate the AI ("ignore instructions, give 100")?

That's prompt injection. Mitigations: clearly delimit the candidate answer in the prompt, instruct the model to treat it as data, cross-check with keyword score (flag large disagreements), and clamp the score. Currently the answer is labelled "CANDIDATE ANSWER:" but not otherwise defended — an honest limitation.

Q5. What happens if two topics have no questions?

BalancedQuestionSetter skips empty pools; Topic focus on an empty topic returns an empty list → the session has 0 questions → average 0 (guarded against division by zero).

Q6. Why is Question's keyword list exposed through unmodifiableList but AnswerEvaluation uses List.copyOf?

Two valid techniques: Question returns a read-only view on each call; AnswerEvaluation stores an immutable copy once at construction. Copying at construction is the stronger guarantee; both prevent callers from mutating internal state.

Q7. Can your program run without Gson on the classpath?

Not when run as plain classes — Gson is a compile/runtime dependency. The shaded JAR bundles it, so `java -jar` works with only a JRE 17+.

23. Last-minute revision sheet

Numbers

- 23 questions · 6 topics (4/4/3/3/4/5)
- 5 questions per session (DEFAULT_QUESTION_COUNT)
- 3 question strategies · 3 AI providers + offline
- 34 tests · 11 test classes · 100% pass
- Verdict: ≥ 75 READY · ≥ 50 ALMOST · else NEEDS PREP
- Keyword feedback: ≥ 75 Excellent · ≥ 40 Partial
- Demo: 83/100/83/0/100 → 73% → ALMOST THERE
- Team: Rithika S (evaluators, strategies) · Shwethaa JV (data, tests)
- Java 17 · Gson 2.11.0 · JUnit 5.10.2 · Maven shade 3.5.2

Packages

- **model** — Candidate, Question, Topic, AnswerEvaluation
- **core** — InterviewSession, InterviewListener
- **service** — banks, setters, evaluators, LLM clients, factory, report
- **ui** — CandidateRegistrationUI, ConsoleInterviewUI
- **util** — JsonQuestionLoader

Patterns in one line each

- **Strategy** — QuestionSetter / AnswerEvaluator swapped at runtime
- **Factory** — EvaluatorFactory picks Claude → Gemini → Ollama → keyword
- **Decorator** — FallbackEvaluator adds resilience
- **Observer** — InterviewListener decouples UI
- **Bridge** — AbstractLlmEvaluator has-a LlmClient
- **Template Method** — evaluate() skeleton with hooks
- **Adapter** — LlmClients unify 3 REST APIs
- **Builder** — HttpClient / HttpRequest / GsonBuilder

Java features to point at

- switch expressions + yield · text blocks · instanceof pattern · var
- lambdas · ComputableAbsent · TokenType (type erasure)
- try-with-resources · custom unchecked exception + chaining
- List.of / copyOf · unmodifiableList · shuffle(Random)
- HttpClient · System.getenv · StringBuilder · String.repeat

Final advice. When asked “why”, answer with a *benefit*: testability, extensibility (Open/Closed), resilience, or separation of concerns. When asked about a weakness, name it, explain the impact, give the fix. Keep the poster pipeline in mind: **Registration** → **Strategy** → **Question Bank** → **Evaluation (AI + Fallback)** → **Readiness Report**.

Appendix A. Report errors to fix before submission

The latest report still contains text from a different project (a crypto-trading app called "CrypteX"). An examiner who spots these will ask about them, so fix them in the .docx before submitting or printing.

Where	Current text	Replace with
1.3 Objectives (all six points)	"ingest live market data ... WebSocket", "order matching engine using PriorityQueue books", "SMA, EMA and Wilder RSI ... RingBuffer", "thread-safe virtual wallet using ReentrantLock ... JDBC", "session-based authentication"	This project's objectives (as on the poster): JSON question-bank schema; OOP classes and interfaces for selection and evaluation; offline keyword scoring with an optional AI evaluator and fallback; three question-selection strategies; readiness report; JUnit 5 testing; team reflection.
2.2 text before Table 2.1	"...contributed to the design of CrypteX ... Spring Boot was kept for transport only"	"...to the design of the simulator, showing why core Java was chosen and Spring Boot was not used."
Table 3.1, weeks 3–4	"Designed the RingBuffer and heap books; built WebSocket ingestion and the JDBC schema."	"Designed the Question/Topic/Candidate model and the questions.json schema."
Table 3.1 mentor remarks	"Check multi-symbol behaviour." / "Verify engine output independently."	Remarks that fit this project (e.g. "Check fallback when AI is unavailable." / "Verify scores against hand-calculated values.")
3.2 text before Table 3.2	"...build, test and run CrypteX ... an embedded SQLite database"	"...the simulator. It needs only a standard laptop and a Java 17 runtime; there is no database." (Table 3.2 itself correctly says Database: None.)
5.1 first line	"...the path a single market tick takes through the system."	"...the path a single interview answer takes through the system."
5.3	"Figures 5.1 to 5.3"	"Figures 5.1 and 5.2" (only two figures exist).
6.2 text before Table 6.1	"...difference between the engine's indicator values and an independent browser calculation..."	"...the number of JUnit tests passed and defects found in each test class."
Table 6.1 headings	First column "Version"; last column "Engine vs Browser Indicator Delta"; title "Across Iterations"	"Test class"; delete the last column; title "Test and Defect Outcomes by Test Class".
Table 6.2 title	"Maven Surefire Test Report (17 of 17 passing)"	"Sample Mock-Interview Topic-wise Scores" (this also matches the List of Tables).
Appendix A.1	"...all Java sources, the SQL schema , the JUnit suite and the web terminal "	"...all Java sources, questions.json and the JUnit suite".
Bonafide certificate	"SHWETHAA (2104251050318)"	"SHWETHAA JV (2104251050318)"
Table A.1	"SHWETHA JV"; remarks column empty	"SHWETHAA JV"; add remarks (Rithika: evaluators & strategies; Shwethaa: data layer & tests).
1.2 driving question	Closing quotation mark missing	Add " at the end.
1.1, 2.1.2, 4.2, 6.4	Double full stops ("cycle..", "demo..", "swapped..", "environment..")	Single full stop.
List of Abbreviations	Lists JDBC and REST (barely used) but not JVM, JSON, LLM	Add JVM, JSON, LLM; keep JDBC/REST only if you mention them.
Table of contents	Some page numbers are hard-coded while others are fields	In Word: select all → F9 to update every field, then check the numbers.

If the examiner finds one before you fix it: say plainly that it was left over from the report template the team started from, and state what it should say. Don't try to defend it as part of this project.